

Computer Vision

CSCI 4420

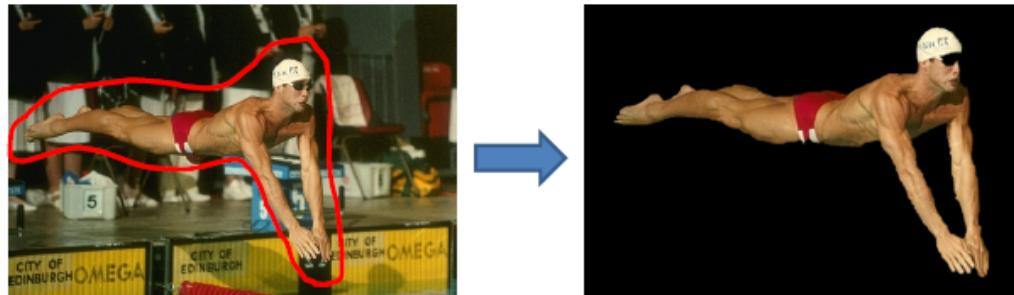
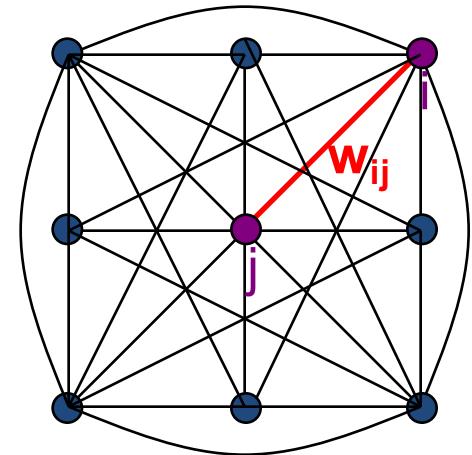
Instructor: Faisal Qureshi

Slides Credit: These slides borrow heavily from on-line resources, particularly from other similar computer vision courses at Brown University, the University of Washington and Stanford.

Segmentation

Today's class

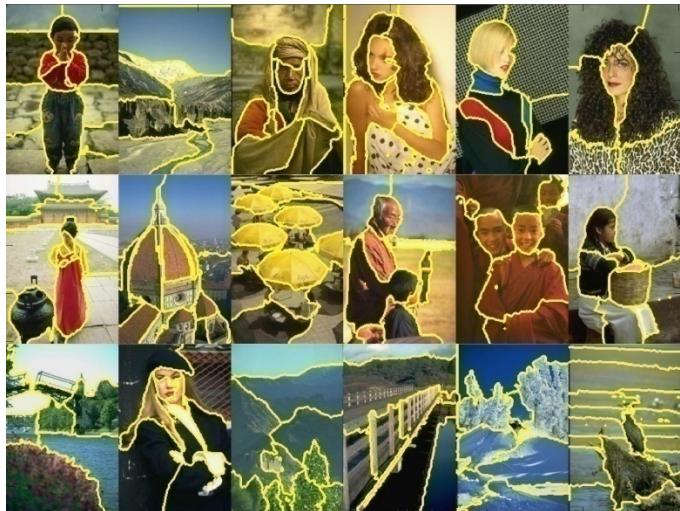
- Segmentation and Grouping
- Inspiration from human perception
 - Gestalt properties
- MRFs
- Segmentation with Graph Cuts



Grouping in vision

- Goals:
 - Gather features that belong together
 - Obtain an intermediate representation that compactly describes key image or video parts

Examples of grouping in vision



[Figure by J. Shi]

Determine image regions

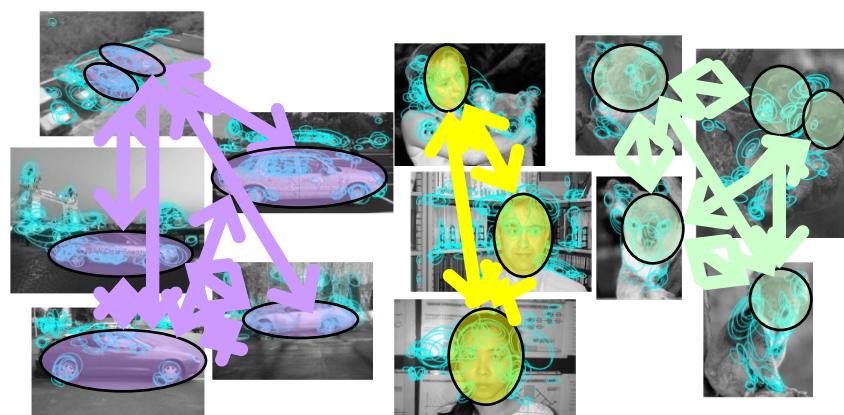


Group video frames into shots



[Figure by Wang & Suter]

Figure-ground



[Figure by Grauman & Darrell]

Object-level grouping

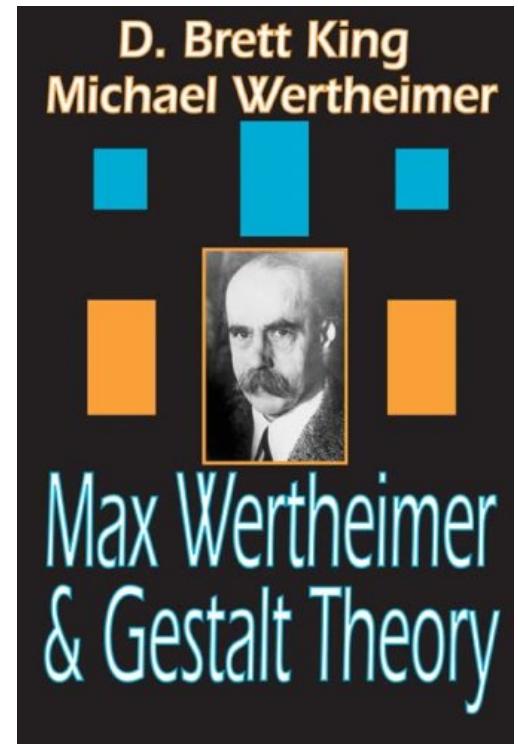
Grouping in vision

- Goals:
 - Gather features that belong together
 - Obtain an intermediate representation that compactly describes key image (video) parts
- Top down vs. bottom up segmentation
 - Top down: pixels belong together because they are from the same object
 - Bottom up: pixels belong together because they look similar
- Hard to measure success
 - What is interesting depends on the app.

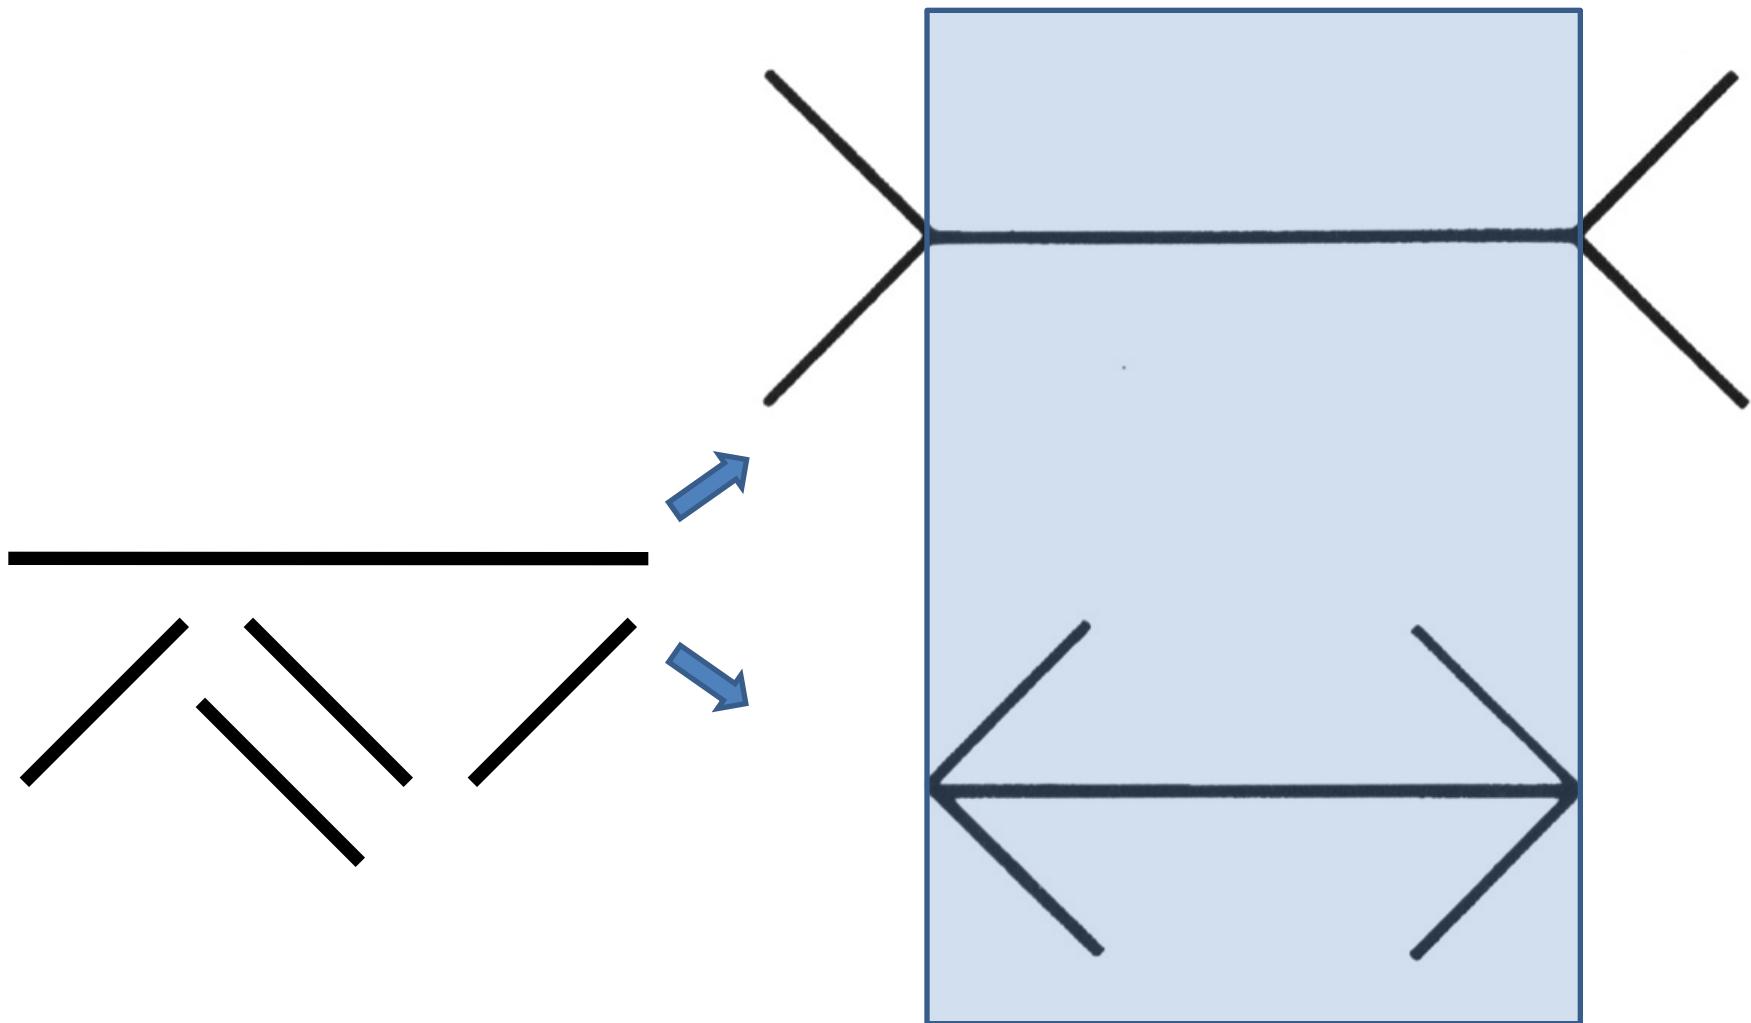
What things should be grouped?
What cues indicate groups?

Gestalt psychology or Gestaltism

- German: *Gestalt* - "form" or "whole"
- Berlin School, early 20th century
 - Kurt Koffka, Max Wertheimer, and Wolfgang Köhler
- Gestalt: whole or group
 - Whole is greater than sum of its parts
 - Relationships among parts can yield new properties/features
- Psychologists identified series of factors that predispose set of elements to be grouped (by human visual system)

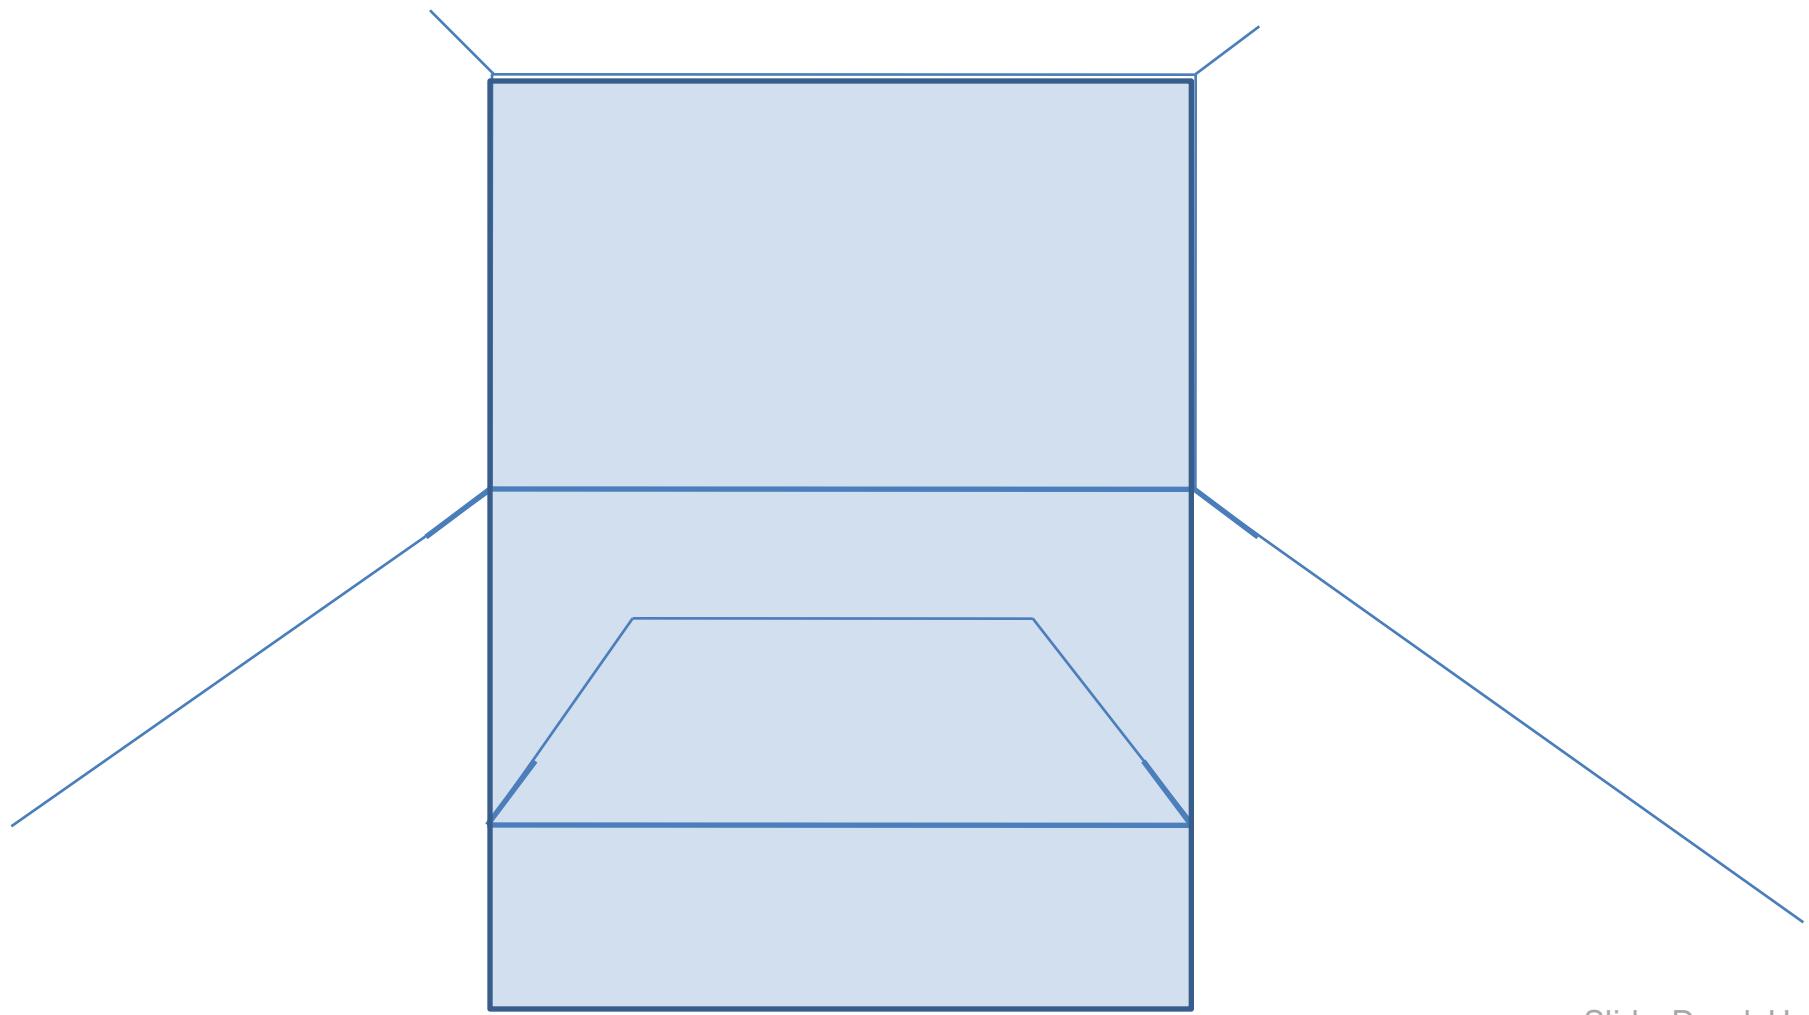


Gestaltism



The Muller-Lyer illusion

We perceive the interpretation, not the senses



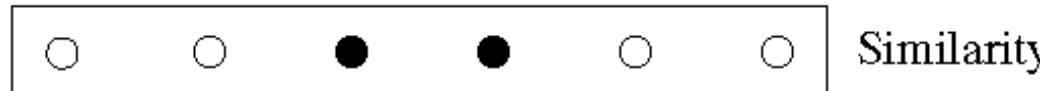
Principles of perceptual organization



Not grouped



Proximity



Similarity



Similarity



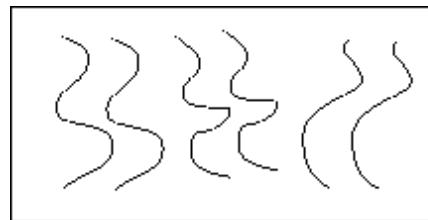
Common Fate



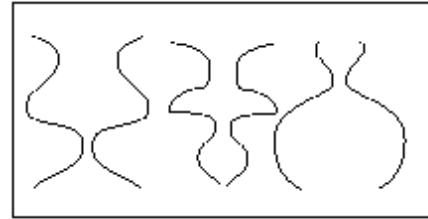
Common Region



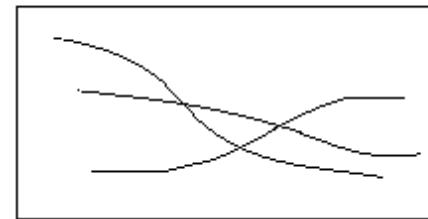
Principles of perceptual organization



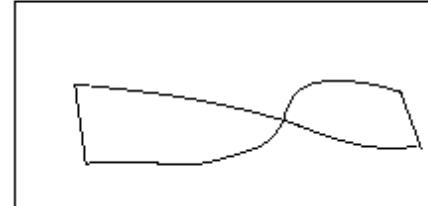
Parallelism



Symmetry



Continuity

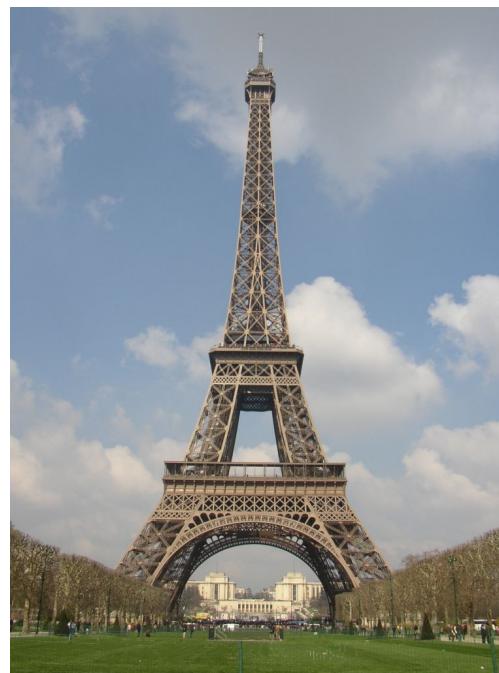


Closure

Similarity



Symmetry



Common fate

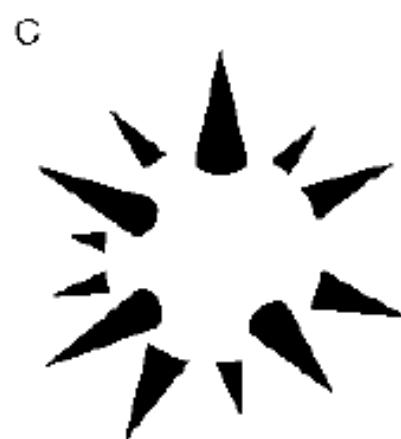
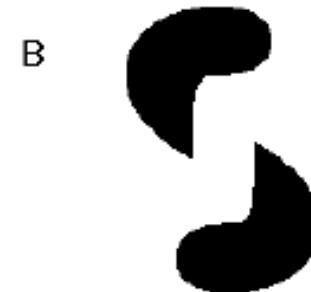
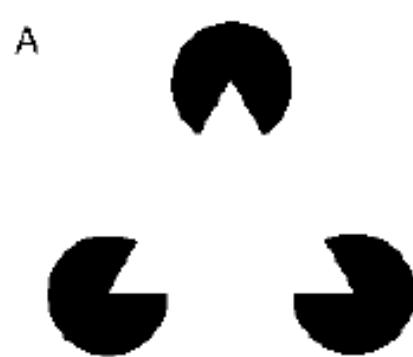


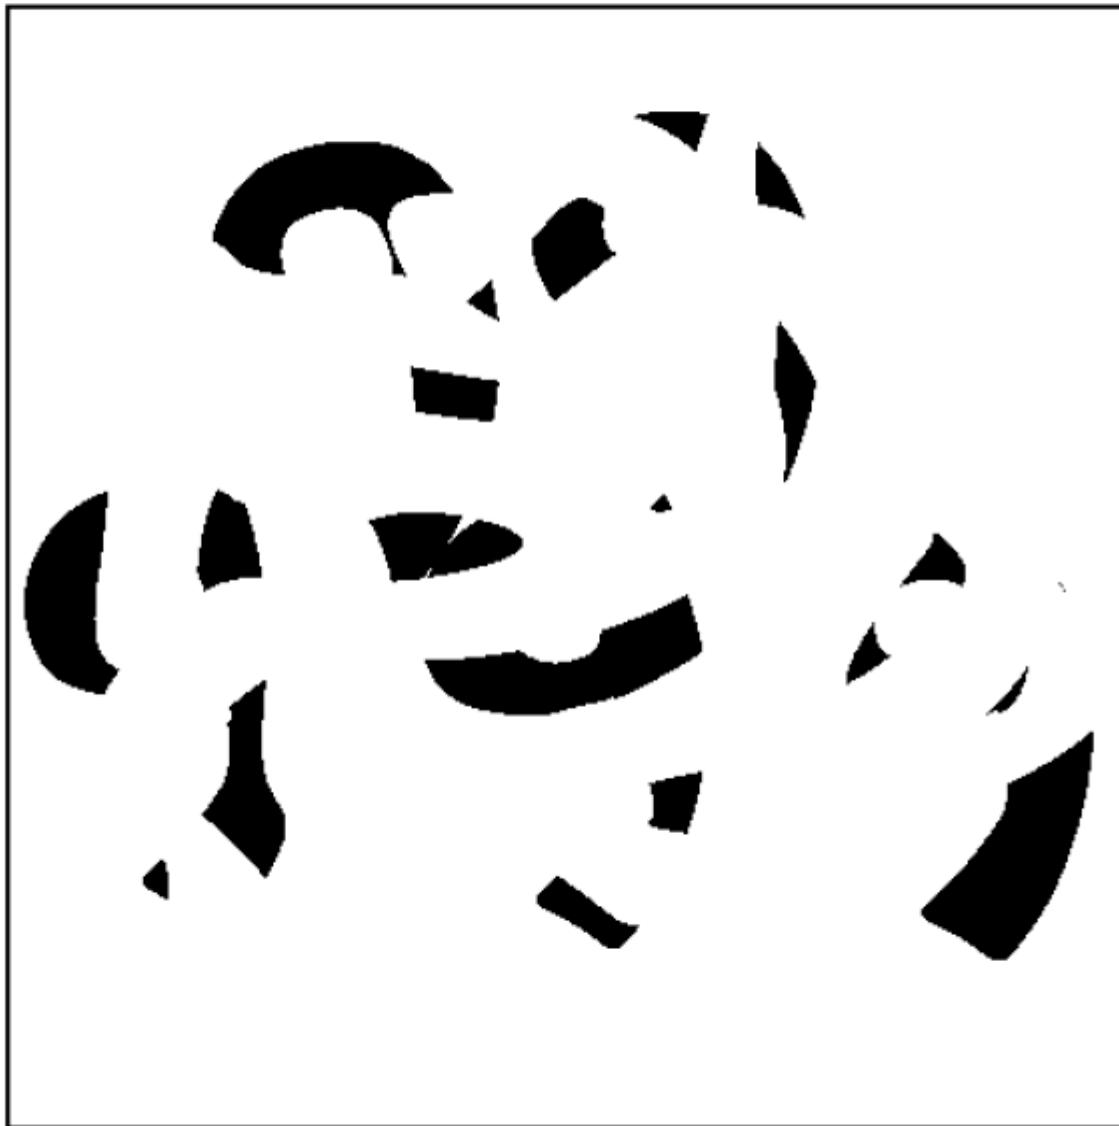
Image credit: Arthus-Bertrand (via F. Durand)

Proximity



Grouping by invisible completion

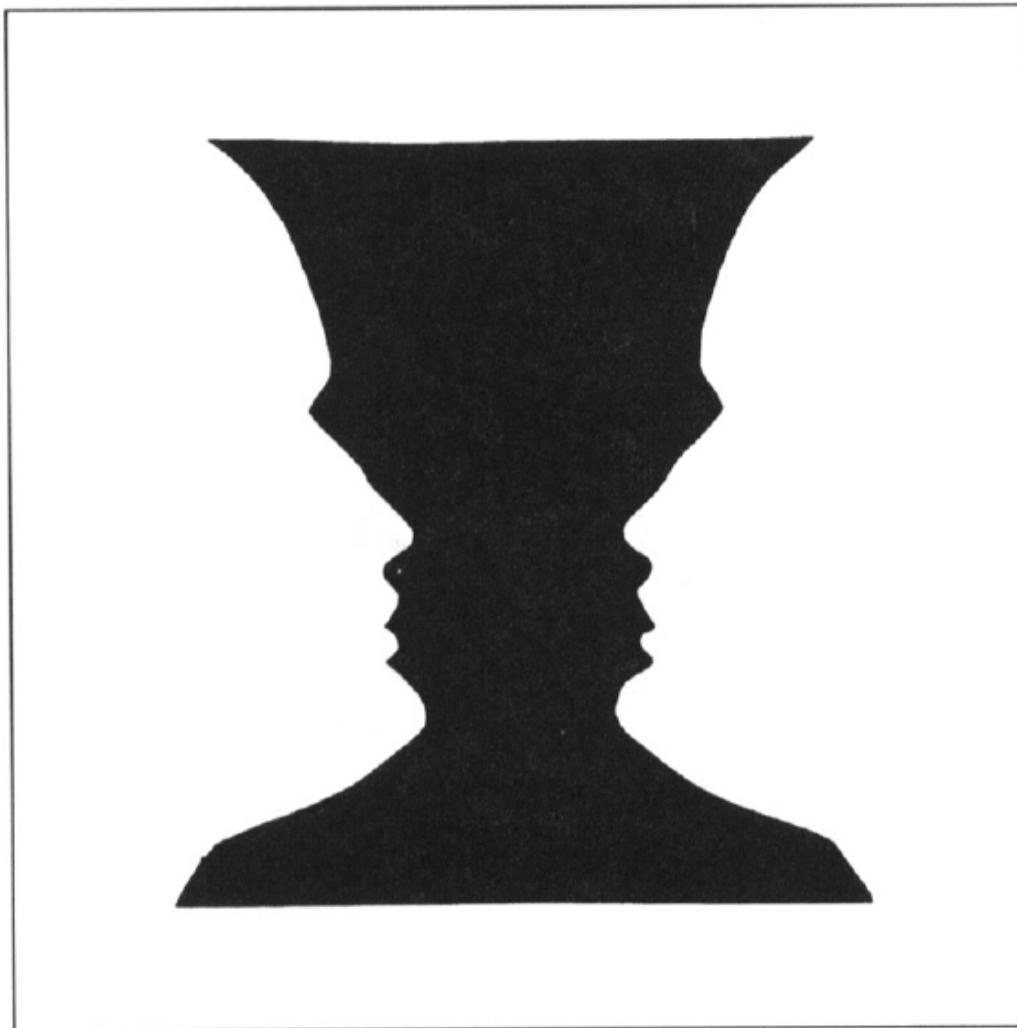




D. Forsyth



Figure-ground



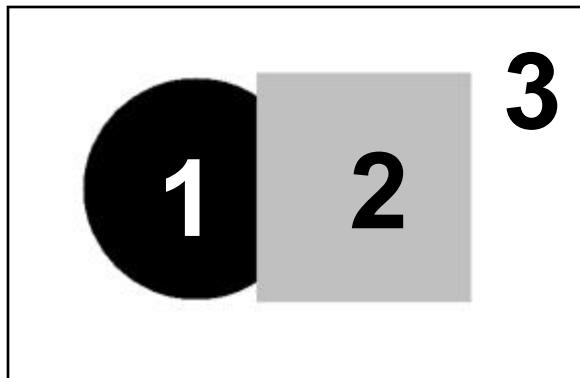
Emergence



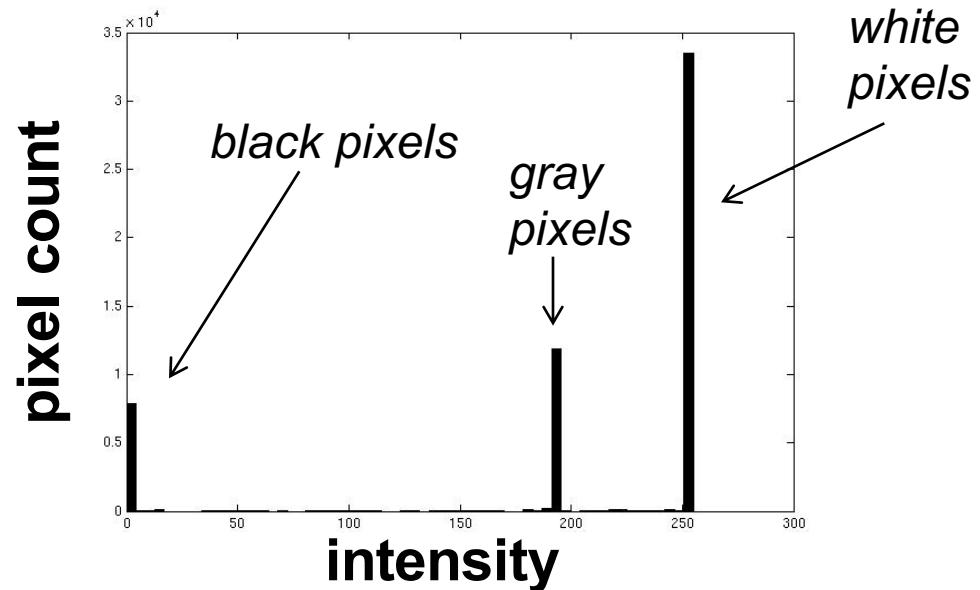
Gestalt cues

- Good intuition and basic principles for grouping
- Basis for many ideas in segmentation and occlusion reasoning
- Some (e.g., symmetry) are difficult to implement in practice

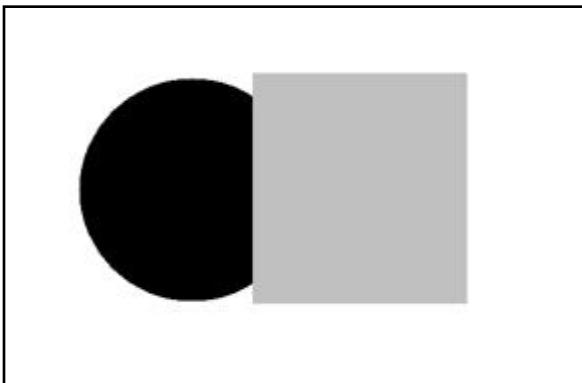
Image segmentation: toy example



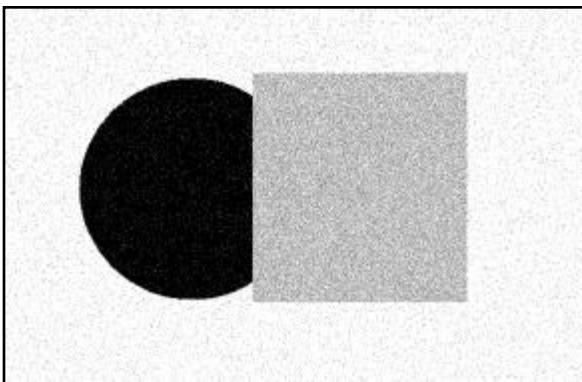
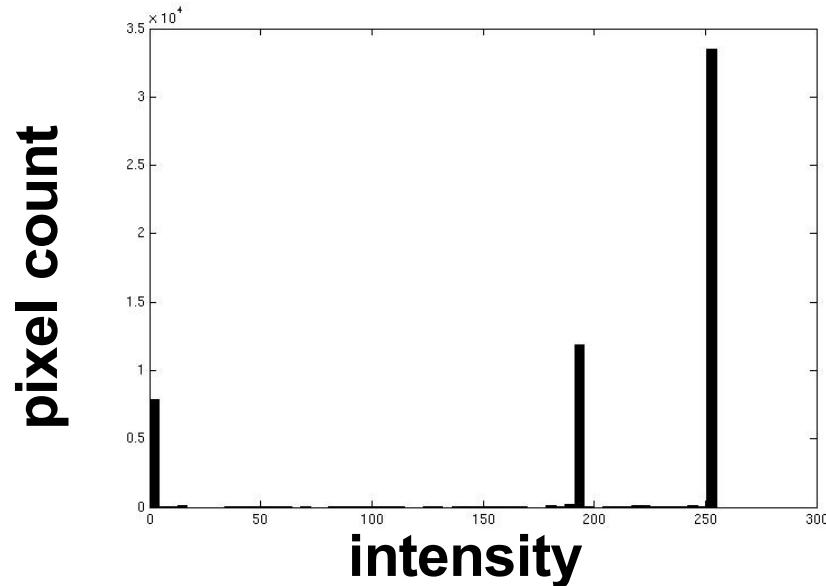
input image



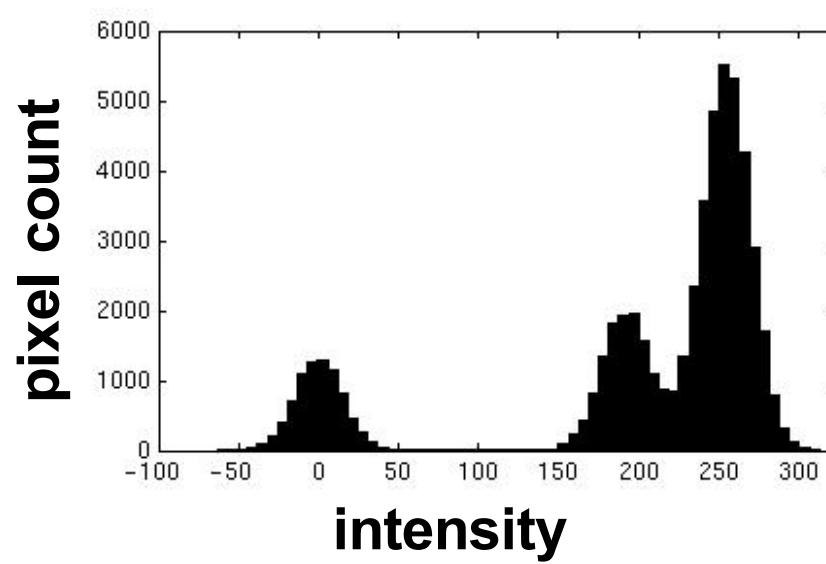
- These intensities define the three groups.
- We could label every pixel in the image according to which of these primary intensities it is.
 - i.e., *segment* the image based on the intensity feature.
- What if the image isn't quite so simple?

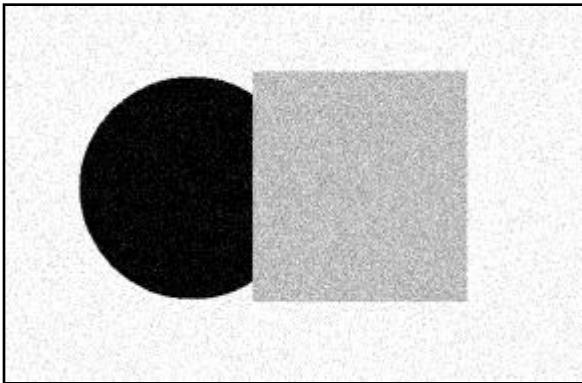


input image

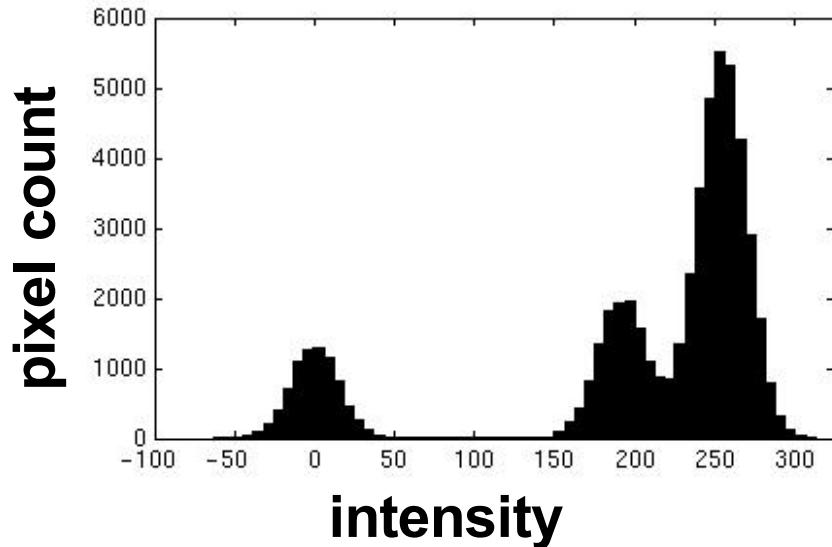


input image

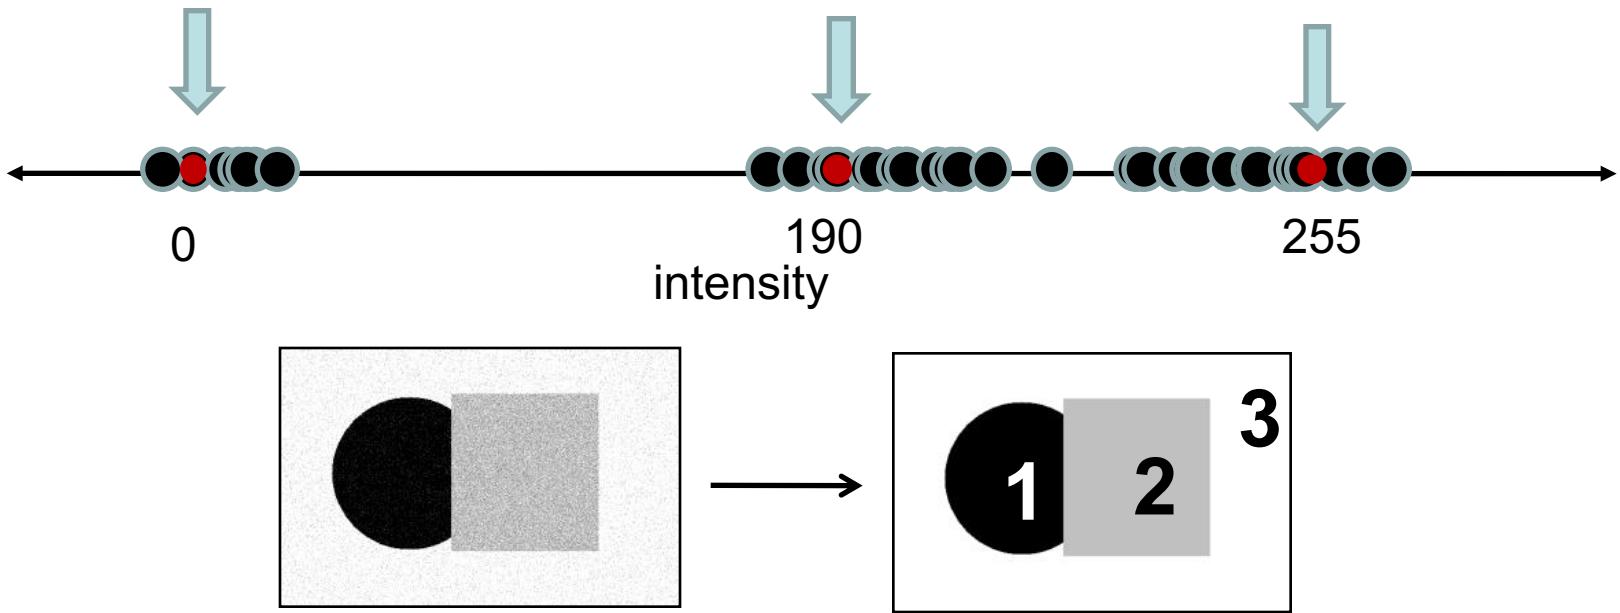




input image



- Now how to determine the three main intensities that define our groups?
- We need to ***cluster***.

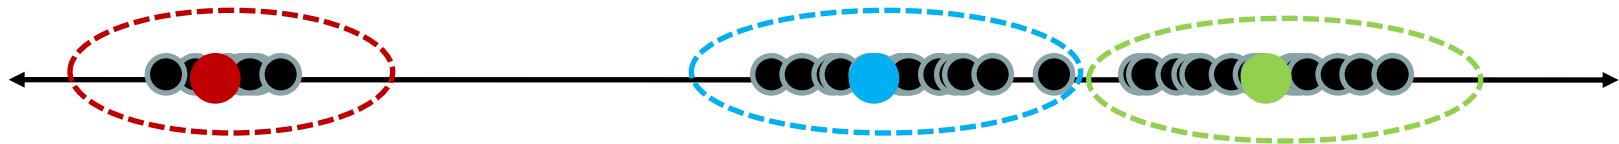


- Goal: choose three “centers” as the **representative** intensities, and label every pixel according to which of these centers it is nearest to.
- Best cluster centers are those that minimize SSD between all points and their nearest cluster center c_i :

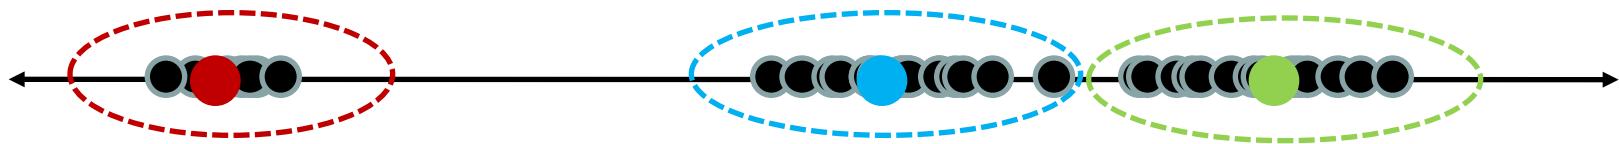
$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

Clustering

- With this objective, it is a “chicken and egg” problem:
 - If we knew the **cluster centers**, we could allocate points to groups by assigning each to its closest center.



- If we knew the **group memberships**, we could get the centers by computing the mean per group.



K-means clustering

- Basic idea: randomly initialize the k cluster centers, and iterate between the two steps we just saw.
 1. Randomly initialize the cluster centers, c_1, \dots, c_K
 2. Given cluster centers, determine points in each cluster
 - For each point p , find the closest c_i . Put p into cluster i
 3. Given points in each cluster, solve for c_i
 - Set c_i to be the mean of points in cluster i
 4. If c_i have changed, repeat Step 2
- 

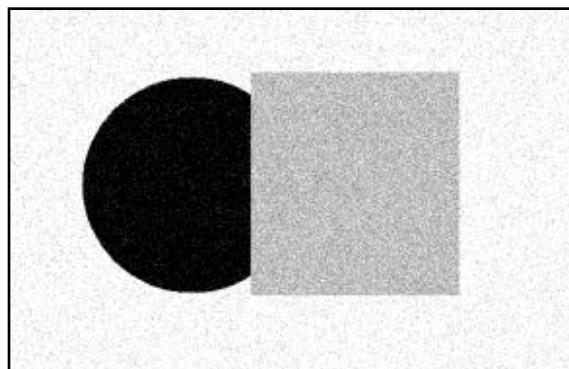
Properties

- Will always converge to some solution
- Can be a “local minimum”
 - does not always find the global minimum of objective function:

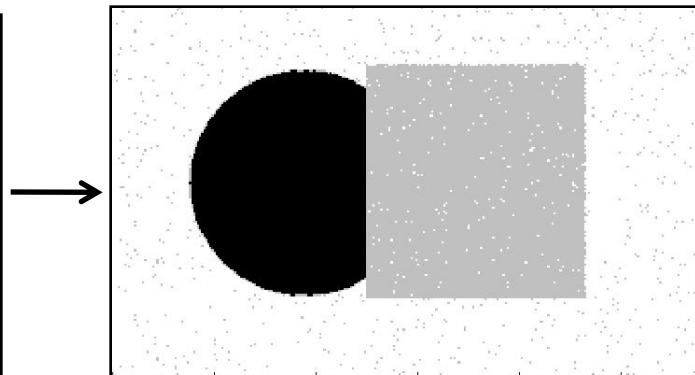
$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

Smoothing out cluster assignments

- Assigning a cluster label per pixel may yield outliers:

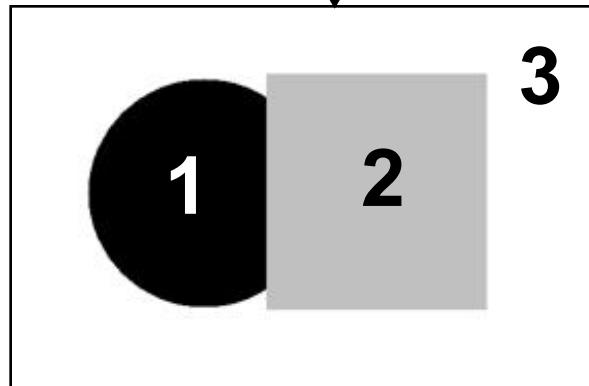


original



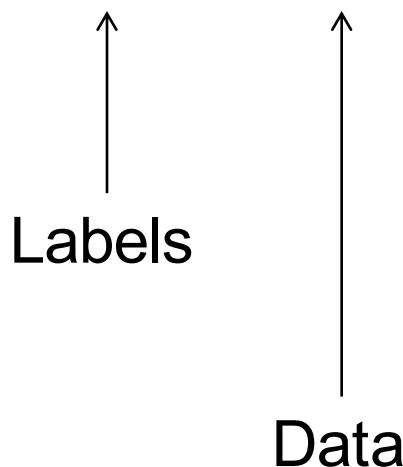
labeled by cluster center's
intensity

- How to ensure they are spatially smooth?



How to Enforce Smoothness?

$$\Pr(w_1 \dots N | x_1 \dots N)$$



How to Enforce Smoothness?

$$\Pr(w_1 \dots N | x_1 \dots N) = \frac{\prod_{n=1}^N \Pr(x_n | w_n) \Pr(w_1 \dots N)}{\Pr(x_1 \dots N)}$$

The diagram illustrates the components of the Bayesian formula. Two vertical arrows point upwards from the labels "Labels" and "Data" to the terms $\Pr(w_1 \dots N)$ and $\Pr(x_1 \dots N)$ respectively in the equation.

How to Enforce Smoothness?

$$\Pr(w_1 \dots N | x_1 \dots N) = \frac{\prod_{n=1}^N \Pr(x_n | w_n)}{\Pr(x_1 \dots N)}$$

Labels

Data

Prior

How about enforcing
a smoothness prior?

How to Enforce Smoothness?

$$\Pr(w_1 \dots N | x_1 \dots N) = \frac{\prod_{n=1}^N \Pr(x_n | w_n)}{\Pr(x_1 \dots N)}$$

Example: Denoise



↑

Prior

How about enforcing
a smoothness prior?

How to Enforce Smoothness?

$$\Pr(w_1 \dots N | x_1 \dots N) = \frac{\prod_{n=1}^N \Pr(x_n | w_n) \Pr(w_1 \dots N)}{\Pr(x_1 \dots N)}$$

Example: Denoise



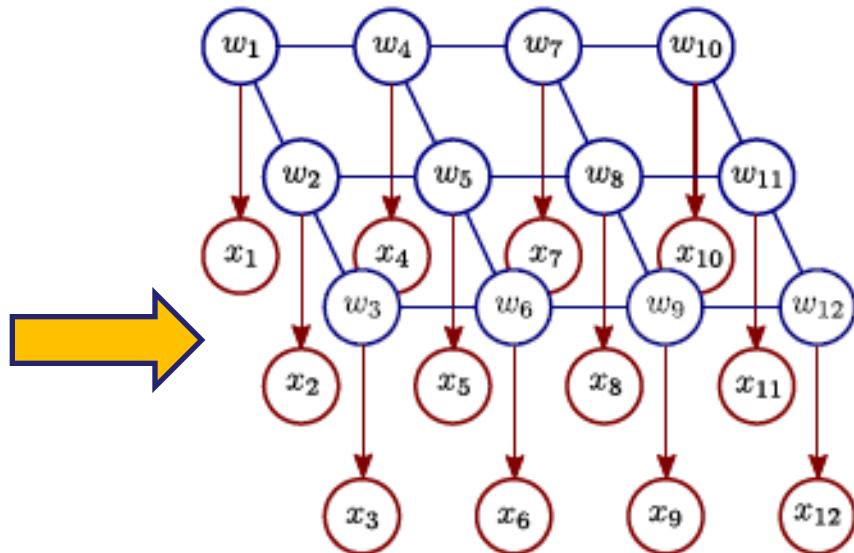
Prior

How about enforcing
a smoothness prior?

How to Enforce Smoothness?

$$\Pr(w_1 \dots N | x_1 \dots N) = \frac{\prod_{n=1}^N \Pr(x_n | w_n) \Pr(w_1 \dots N)}{\Pr(x_1 \dots N)}$$

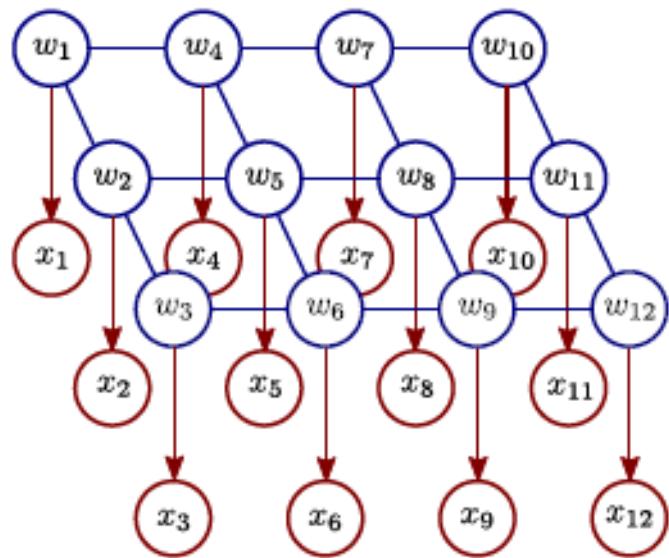
Example: Denoise



Copyright© by Simon Prince

How to Enforce Smoothness?

$$\Pr(w_1 \dots N | x_1 \dots N) = \frac{\prod_{n=1}^N \Pr(x_n | w_n) \Pr(w_1 \dots N)}{\Pr(x_1 \dots N)}$$

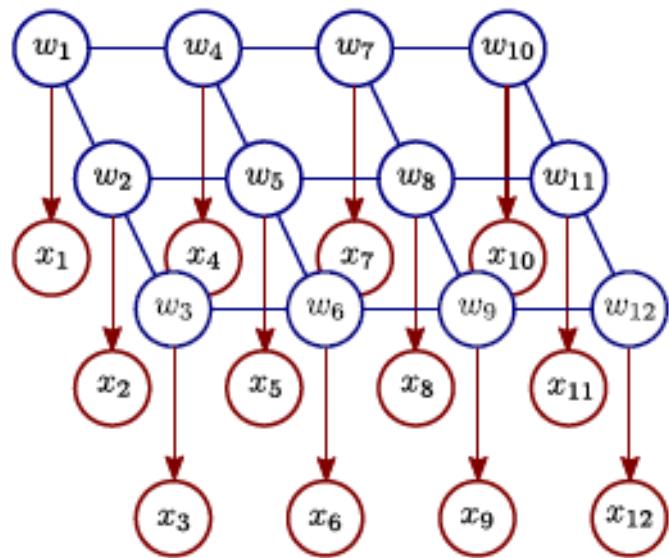


Prior

$$\Pr(w_1 \dots N) = \frac{1}{Z} \exp\left(- \sum_{(m,n) \in C} \varphi(w_m, w_n)\right)$$

How to Enforce Smoothness?

$$\Pr(w_1 \dots N | x_1 \dots N) = \frac{\prod_{n=1}^N \Pr(x_n | w_n) \Pr(w_1 \dots N)}{\Pr(x_1 \dots N)}$$



Prior

$$\Pr(w_1 \dots N) = \frac{1}{Z} \exp \left(- \sum_{(m,n) \in C} \varphi(w_m, w_n) \right)$$

Copyright © by Simon Prince

Neighboring pixels are more likely to have similar labels

Maximum APosteriori Inference

$$\Pr(w_{1\dots N}|x_{1\dots N}) = \frac{\prod_{n=1}^N \Pr(x_n|w_n) \Pr(w_{1\dots N})}{\Pr(x_{1\dots N})}$$

MAP

We aim to find $\{w_n\}$ that maximizes the posterior probability $\Pr(w_{1\dots N}|x_{1\dots N})$.

$$\begin{aligned}\hat{w}_{1\dots N} &= \operatorname{argmax}_{w_{1\dots N}} [\Pr(w_{1\dots N}|x_{1\dots N})] \\ &= \operatorname{argmax}_{w_{1\dots N}} \left[\prod_{n=1}^N \Pr(x_n|w_n) \Pr(w_{1\dots N}) \right] \\ &= \operatorname{argmax}_{w_{1\dots N}} \left[\sum_{n=1}^N \log[\Pr(x_n|w_n)] + \log[\Pr(w_{1\dots N})] \right]\end{aligned}$$

Maximum APosteriori Inference

$$\begin{aligned}\hat{w}_{1\dots N} &= \operatorname{argmax}_{w_{1\dots N}} [Pr(w_{1\dots N} | \mathbf{x}_{1\dots N})] \\ &= \operatorname{argmax}_{w_{1\dots N}} \left[\prod_{n=1}^N Pr(x_n | w_n) Pr(w_{1\dots N}) \right] \\ &= \operatorname{argmax}_{w_{1\dots N}} \left[\sum_{n=1}^N \log[Pr(x_n | w_n)] + \log[Pr(w_{1\dots N})] \right]\end{aligned}$$

Our prior is a Markov Random Field (MRF) with pairwise connections

$$\begin{aligned}\hat{w}_{1\dots N} &= \operatorname{argmax}_{w_{1\dots N}} \left[\sum_{n=1}^N \log[Pr(x_n | w_n)] - \sum_{(m,n) \in C} \psi[w_m, w_n, \theta] \right] \\ &= \operatorname{argmin}_{w_{1\dots N}} \left[\sum_{n=1}^N U_n(w_n) + \sum_{(m,n) \in C} P_{mn}(w_m, w_n) \right],\end{aligned}$$

Maximum APosteriori Inference

$$\begin{aligned}\hat{w}_{1\dots N} &= \underset{w_1\dots w_N}{\operatorname{argmax}} \left[\sum_{n=1}^N \log[Pr(x_n|w_n)] - \sum_{(m,n) \in C} \psi[w_m, w_n, \theta] \right] \\ &= \underset{w_1\dots w_N}{\operatorname{argmin}} \left[\sum_{n=1}^N U_n(w_n) + \sum_{(m,n) \in C} P_{mn}(w_m, w_n) \right],\end{aligned}$$

$U_n(w_n)$ denotes the unary term at pixel n .

Cost of observing the value x_n given w_n .

$P_{mn}(w_m, w_n)$ denotes the pairwise term.

Cost of assigning labels w_n and w_m at neighboring pixels n and m .

Maximum APosteriori Inference

$$\begin{aligned}\hat{w}_{1\dots N} &= \underset{w_{1\dots N}}{\operatorname{argmax}} \left[\sum_{n=1}^N \log[Pr(x_n|w_n)] - \sum_{(m,n) \in C} \psi[w_m, w_n, \theta] \right] \\ &= \underset{w_{1\dots N}}{\operatorname{argmin}} \left[\sum_{n=1}^N U_n(w_n) + \sum_{(m,n) \in C} P_{mn}(w_m, w_n) \right],\end{aligned}$$

$U_n(w_n)$ denotes the unary term at pixel n .

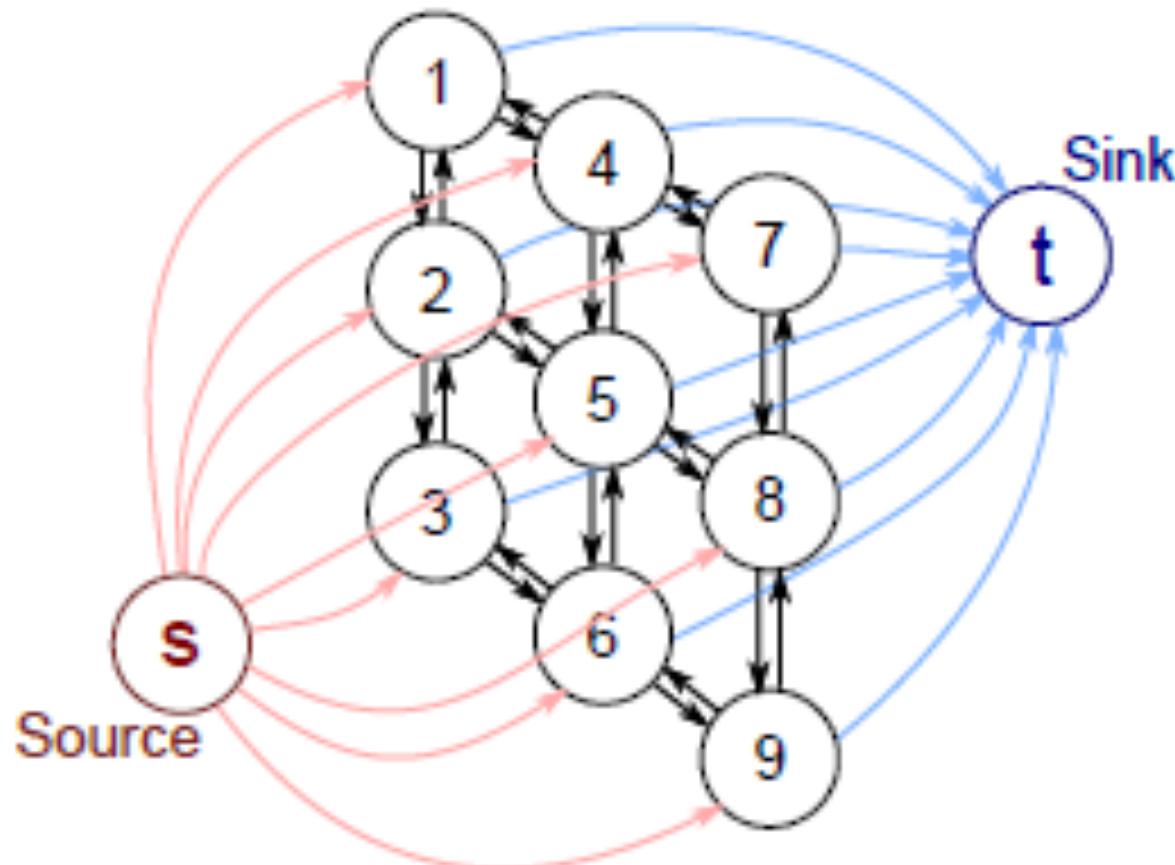
Cost of observing the value x_n given w_n .

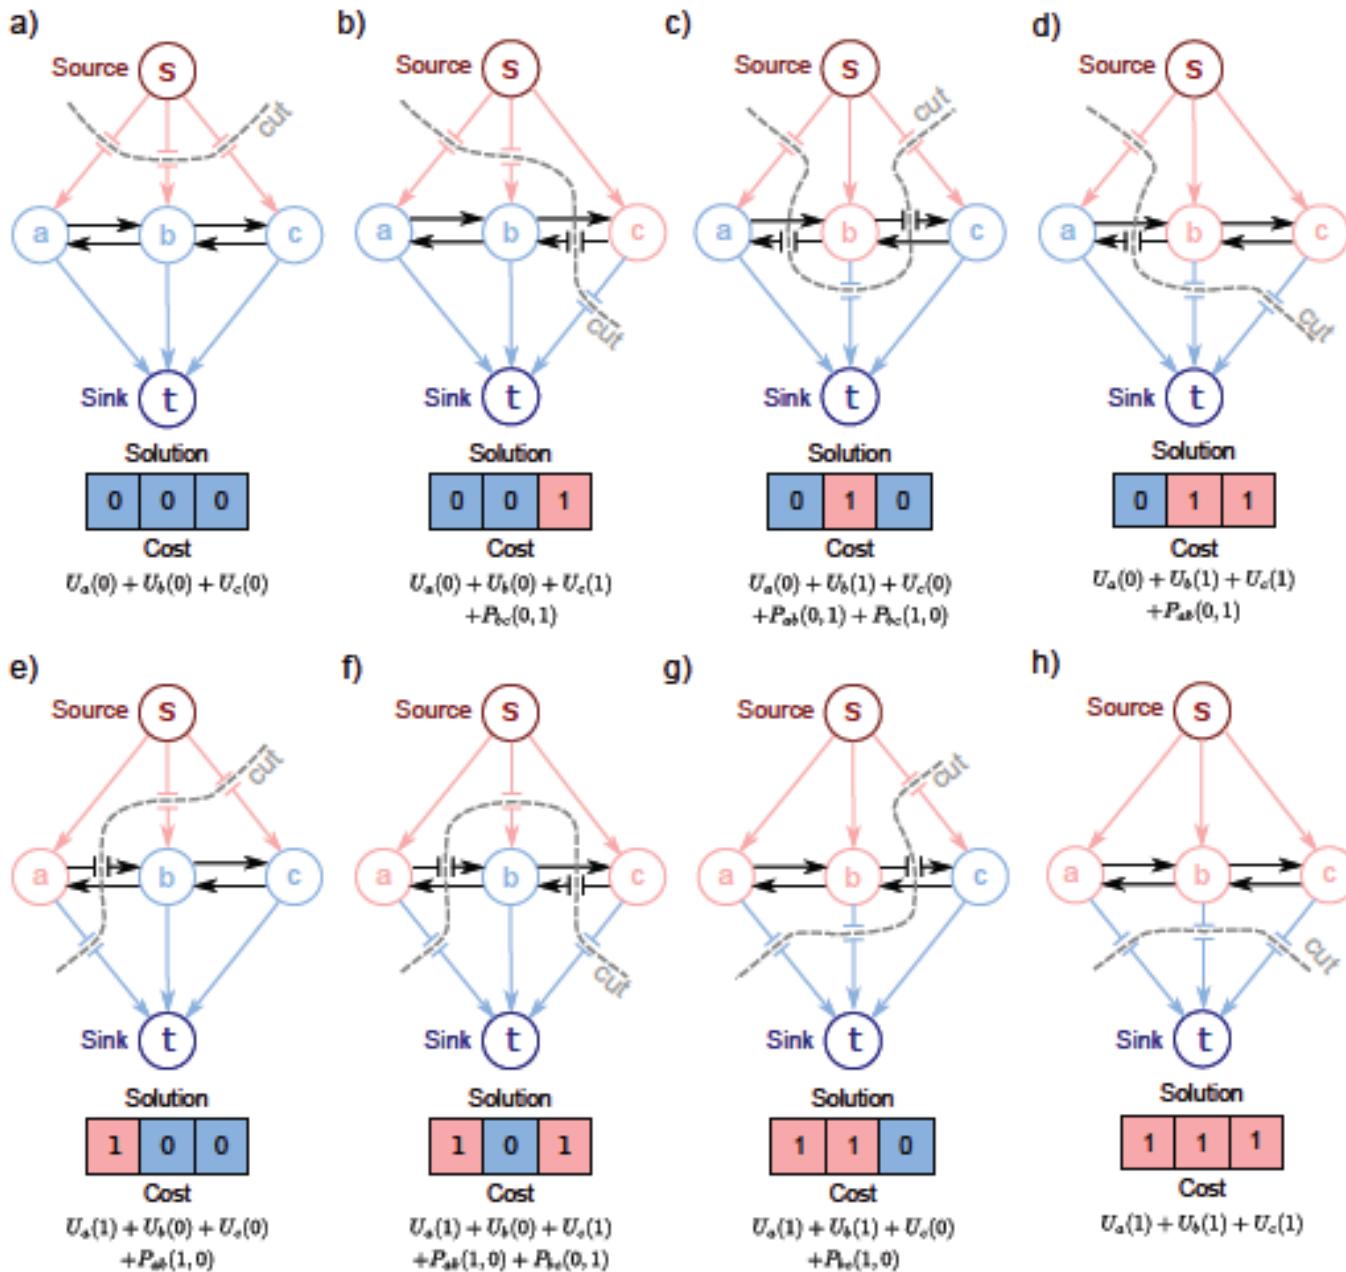
This cost function can be solved using a set of techniques called Graph Cuts.

$P_{mn}(w_m, w_n)$ denotes the pairwise term.

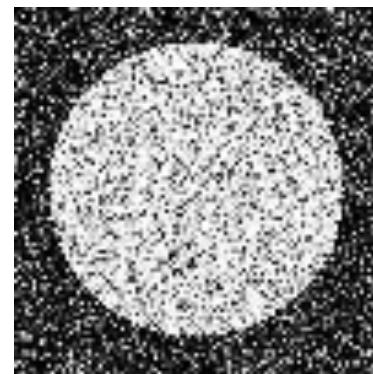
Cost of assigning labels w_n and w_m at neighboring pixels n and m .

Detour: Graph Cuts





Solution



$P(\text{foreground} \mid \text{image})$

Encode dependencies between pixels

Normalizing constant

$$P(\mathbf{y} \mid \theta, \text{data}) = \frac{1}{Z} \prod_{i=1..N} f_1(y_i \mid \theta, \text{data}) \prod_{i,j \in \text{edges}} f_2(y_i, y_j \mid \theta, \text{data})$$

Labels to be predicted Individual predictions Pairwise predictions

Writing Likelihood as an “Energy”

$$P(\mathbf{y} \mid \theta, \text{data}) = \frac{1}{Z} \prod_{i=1..N} p_1(y_i \mid \theta, \text{data}) \prod_{i,j \in \text{edges}} p_2(y_i, y_j \mid \theta, \text{data})$$

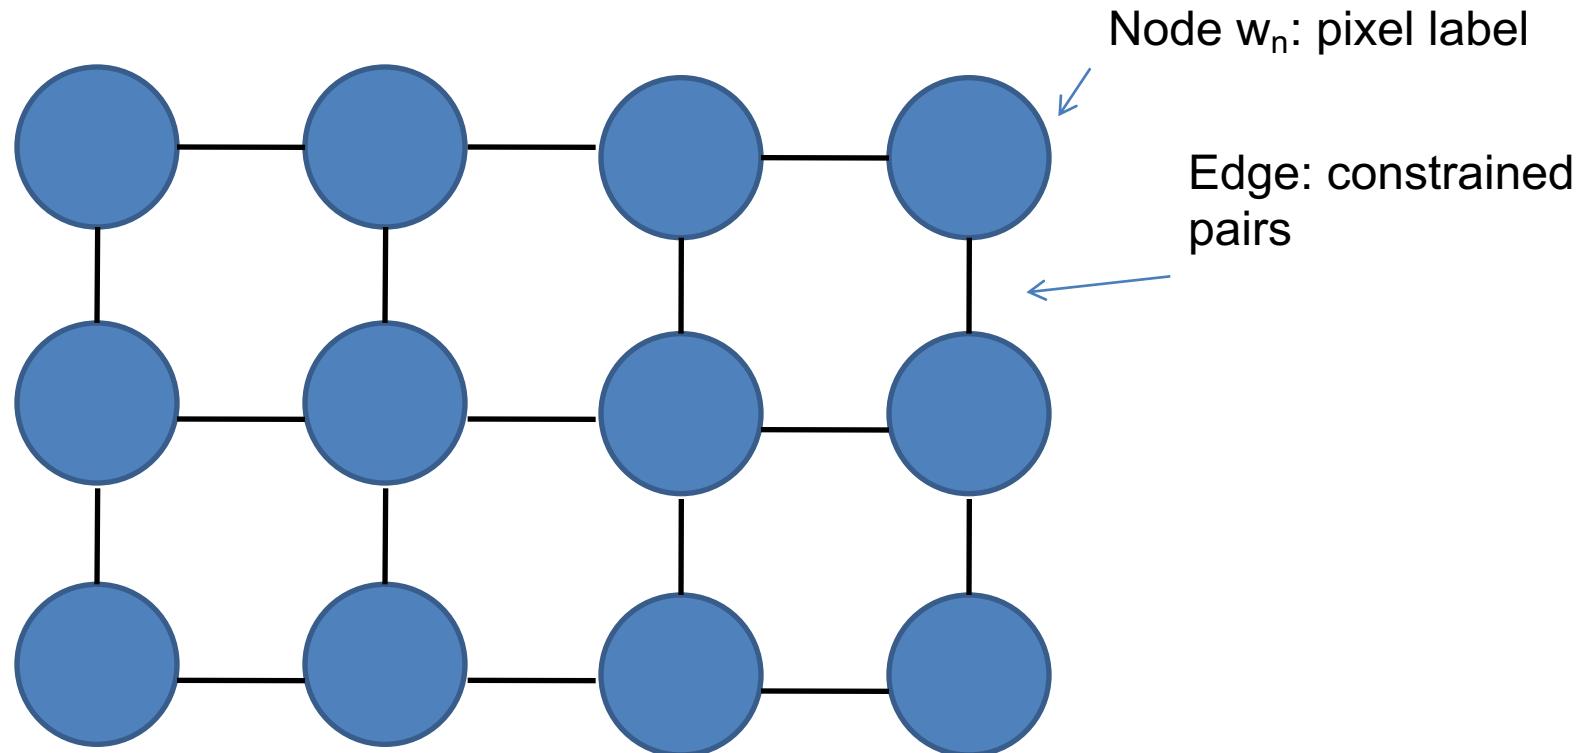


$$\text{Energy}(\mathbf{y} \mid \theta, \text{data}) = \sum_i \psi_1(y_i \mid \theta, \text{data}) + \sum_{i,j \in \text{edges}} \psi_2(y_i, y_j \mid \theta, \text{data})$$

“Cost” of assignment y_i

“Cost” of pairwise assignment y_i, y_j

Markov Random Fields

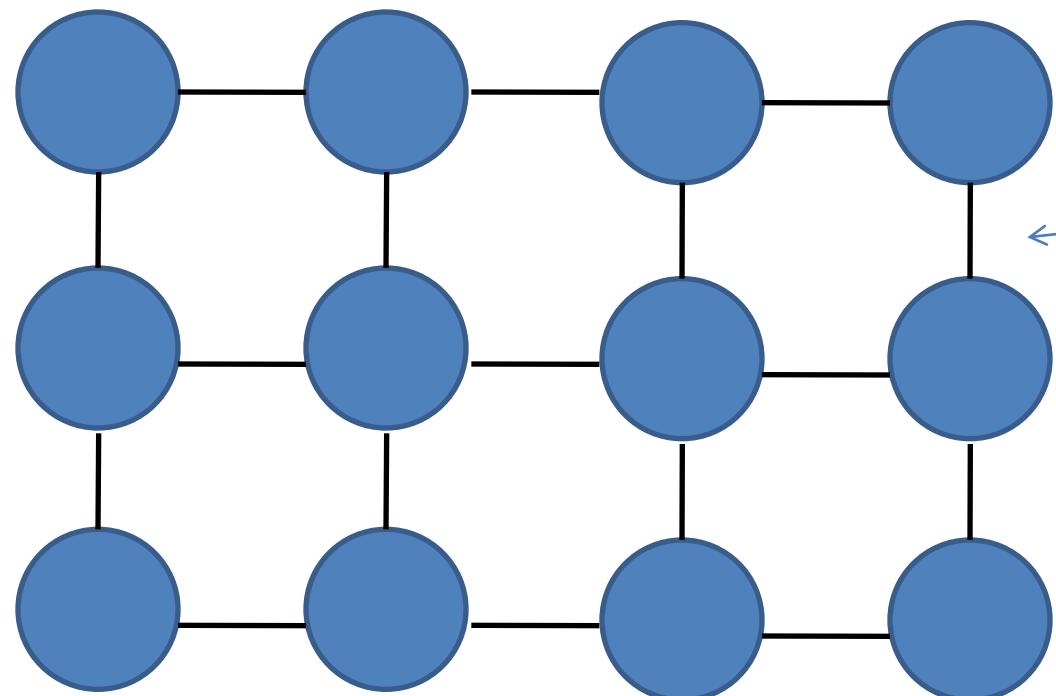


$$Energy(w | \theta, data) = \sum_n U_n(w_n) + \sum_{(m,n) \in \text{edges}} P_{mn}(w_m, w_n)$$

Source: Derek Hoiem

Markov Random Fields

- Example: “label smoothing” grid



Cost to assign a label to
each pixel

Cost to assign a pair of labels to
connected pixels

$$Energy(w | \theta, data) = \sum_n U_n(w_n) + \sum_{(m,n) \in edges} P_{mn}(w_m, w_n)$$

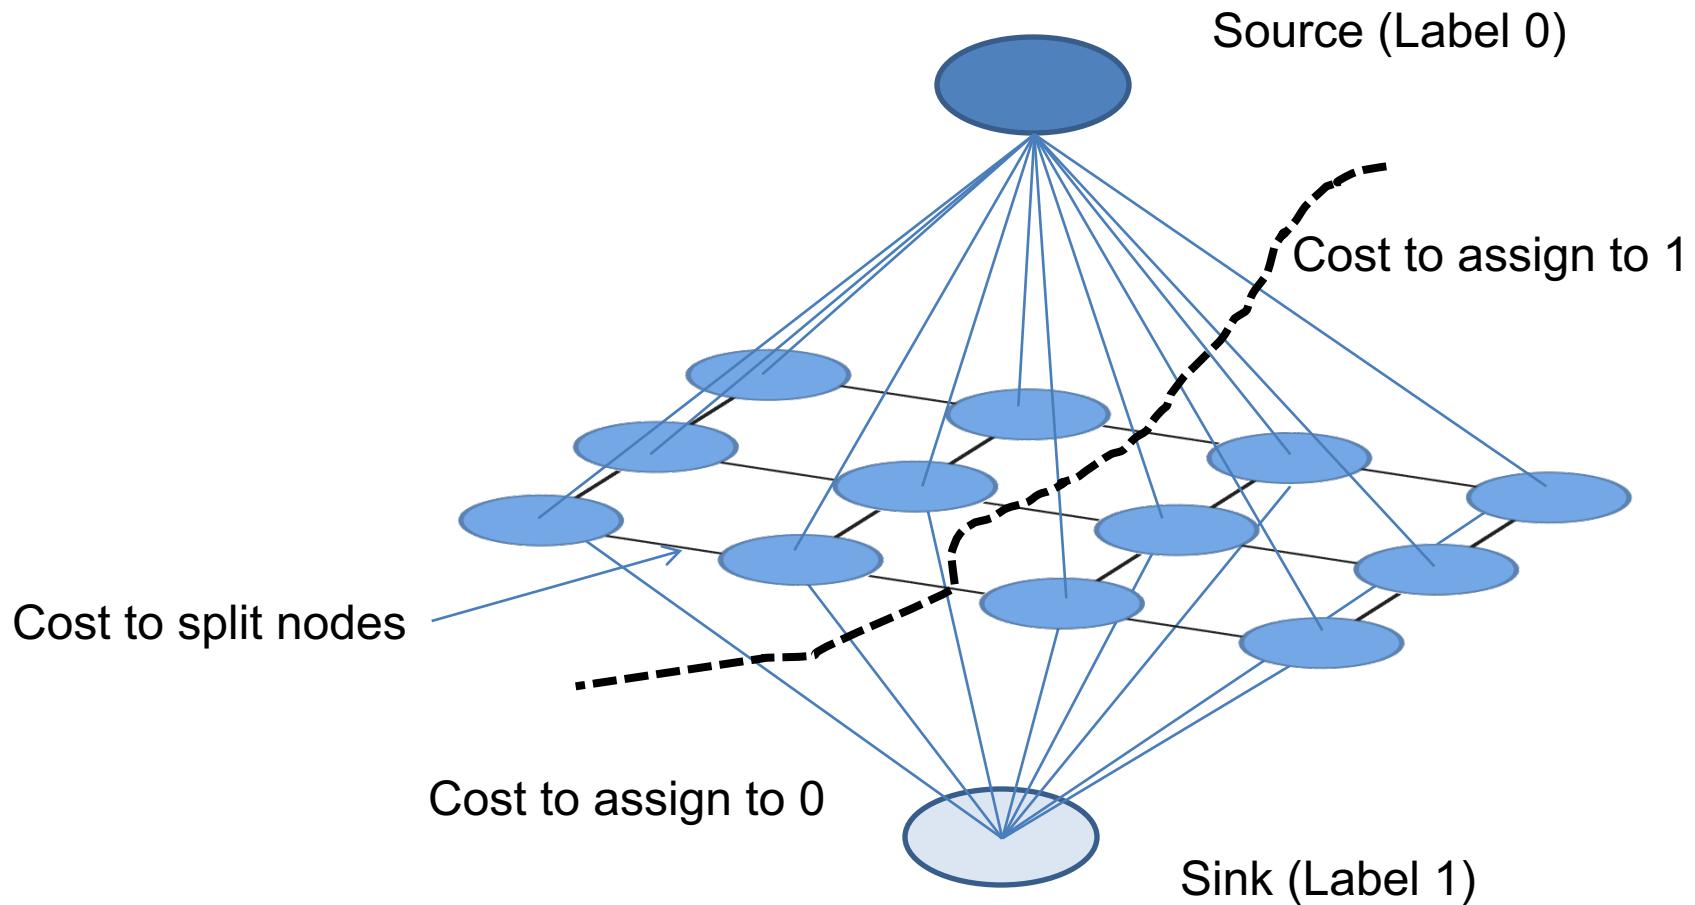
Unary potential

0: $-\log P(w_n = 0 | \text{data})$
1: $-\log P(w_n = 1 | \text{data})$

Pairwise Potential

	0	1
0	0	K
1	K	0

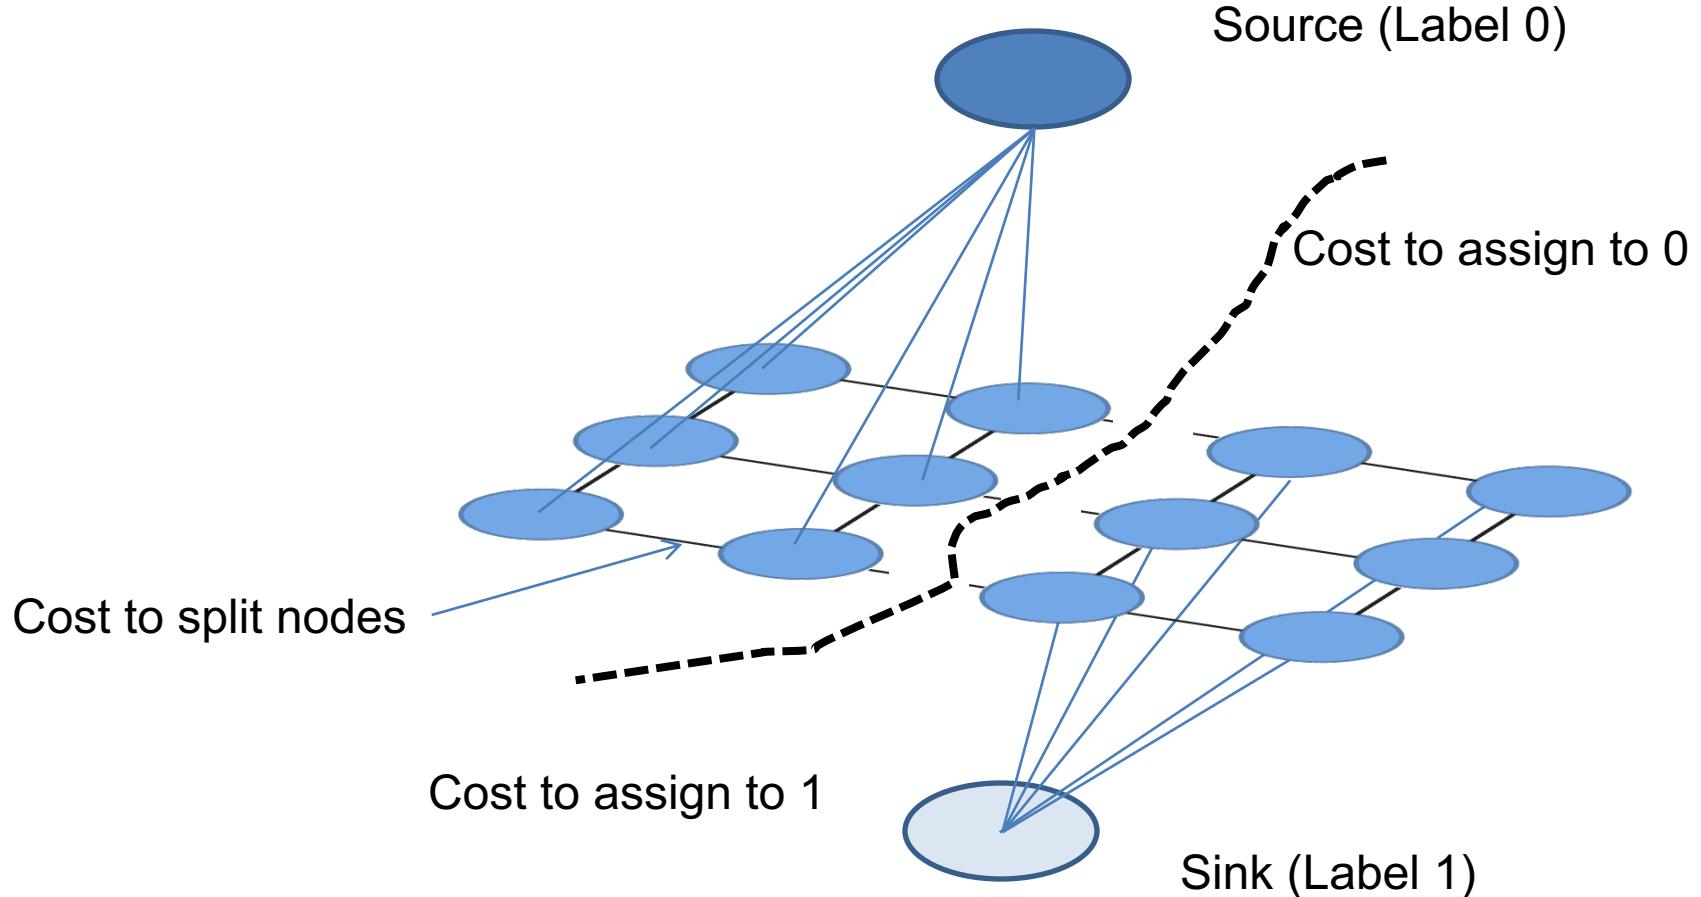
Solving MRFs with graph cuts



$$Energy(w | \theta, data) = \sum_n U_n(w_n) + \sum_{(m,n) \in \text{edges}} P_{mn}(w_m, w_n)$$

Source: Derek Hoiem

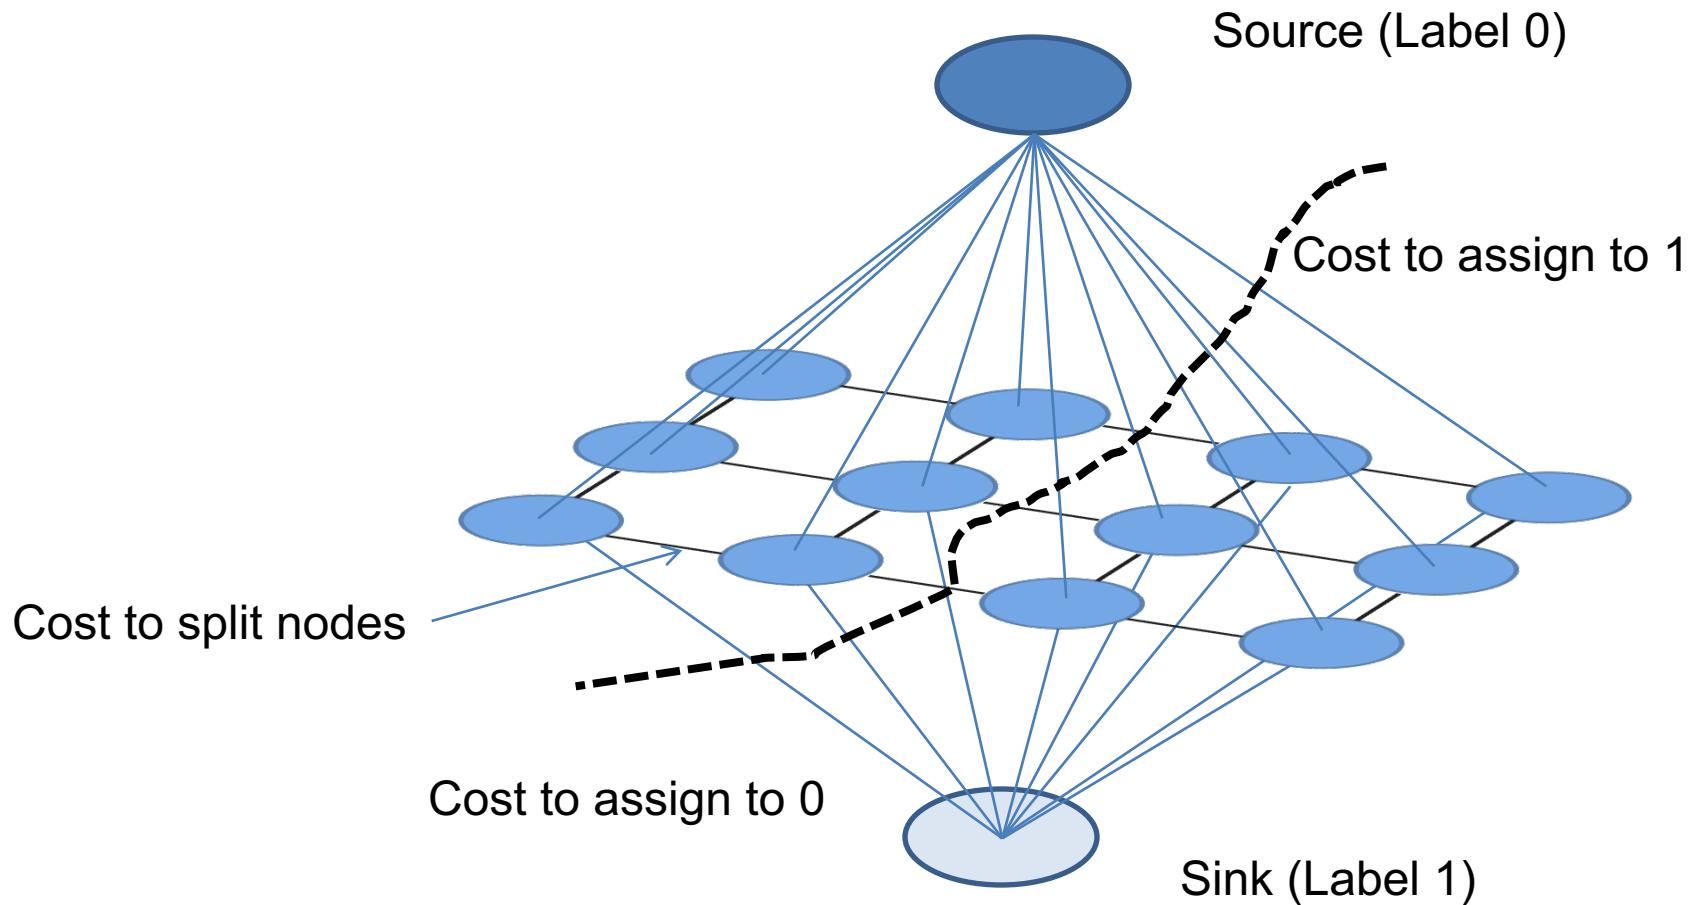
Solving MRFs with graph cuts



$$Energy(w | \theta, data) = \sum_n U_n(w_n) + \sum_{(m,n) \in edges} P_{mn}(w_m, w_n)$$

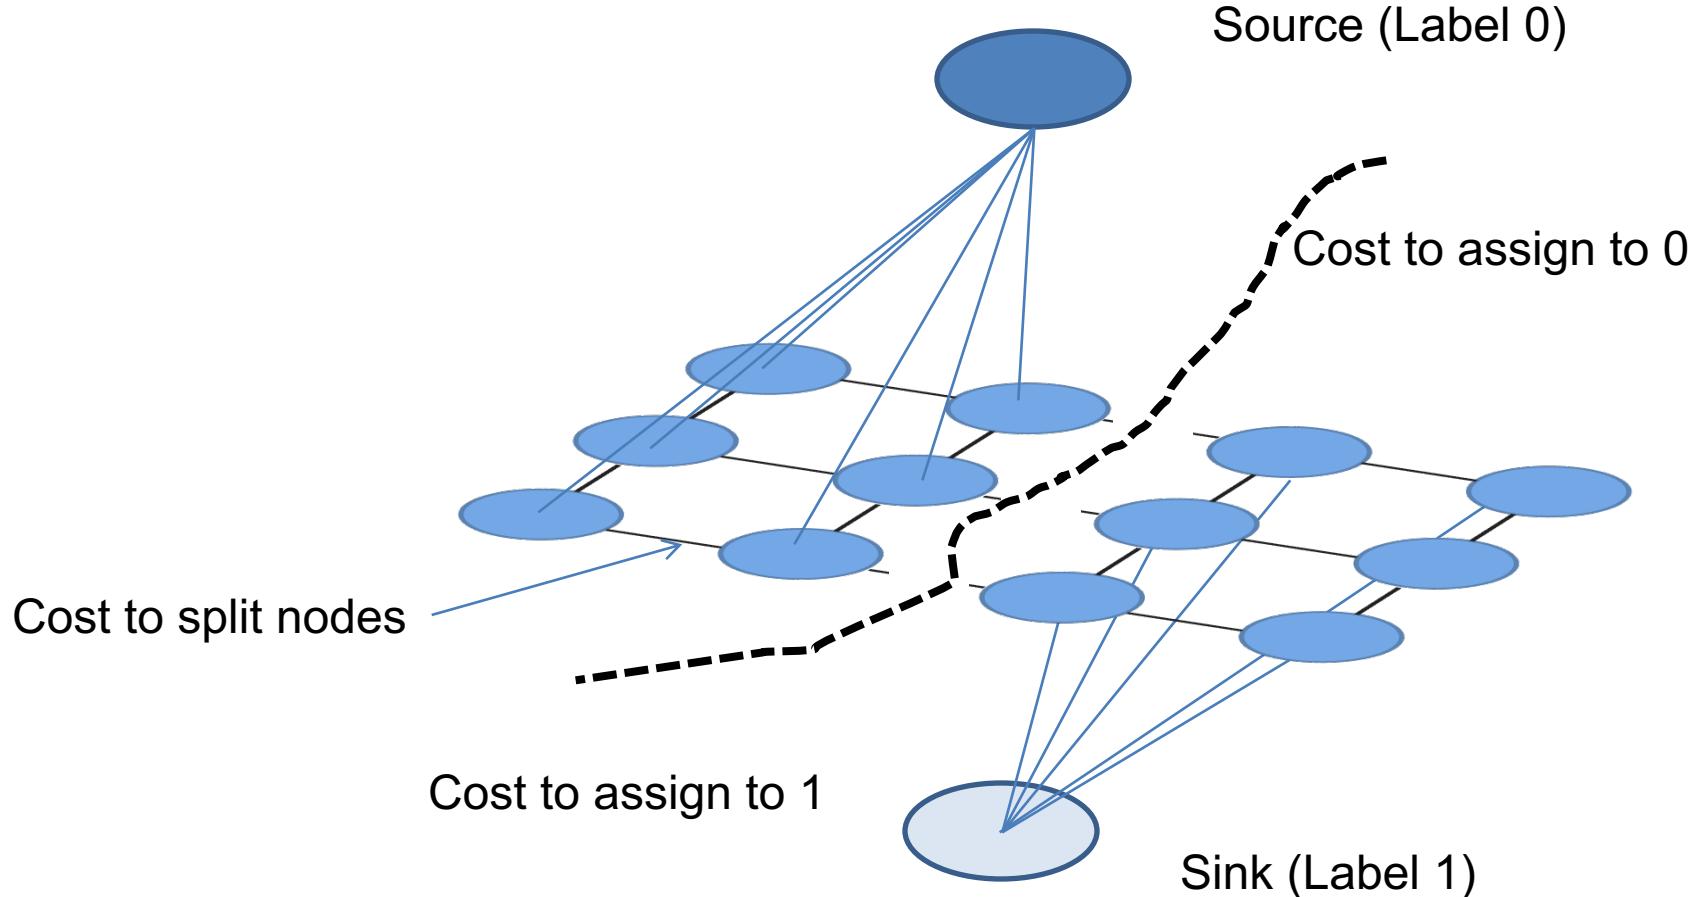
Source: Derek Hoiem

Solving MRFs with graph cuts



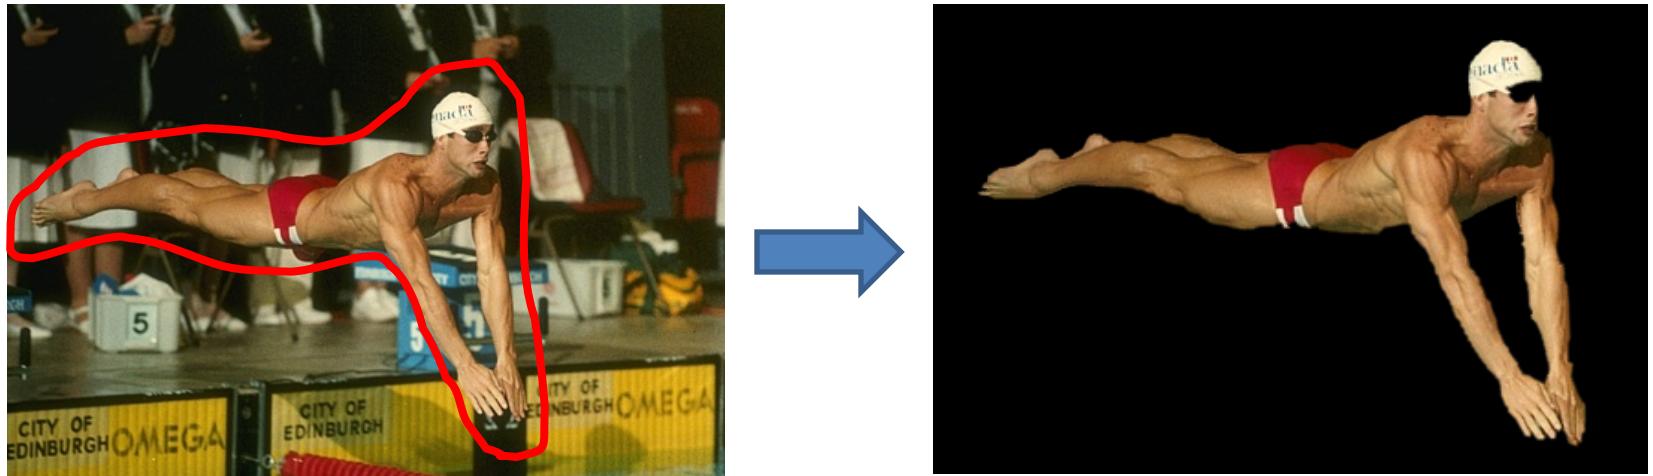
$$Energy(\mathbf{y} \mid \theta, data) = \sum_i \psi_1(y_i \mid \theta, data) + \sum_{i,j \in edges} \psi_2(y_i, y_j \mid \theta, data)$$

Solving MRFs with graph cuts



$$Energy(\mathbf{y} \mid \theta, data) = \sum_i \psi_1(y_i \mid \theta, data) + \sum_{i,j \in edges} \psi_2(y_i, y_j \mid \theta, data)$$

GrabCut segmentation



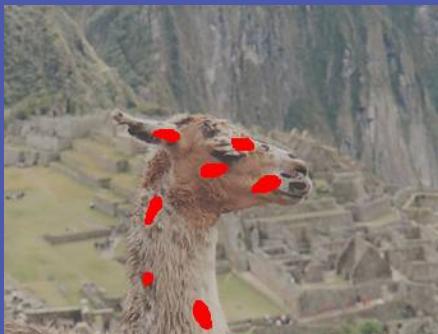
User provides rough indication of foreground region.

Goal: Automatically provide a pixel-level segmentation.

Grab cuts and graph cuts

User
Input

Magic Wand
(198?)



Regions

Intelligent Scissors
Mortensen and Barrett (1995)



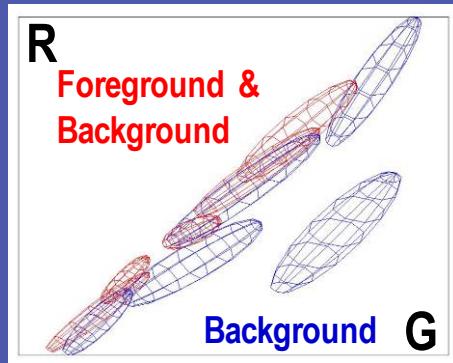
Boundary

GrabCut



Regions & Boundary

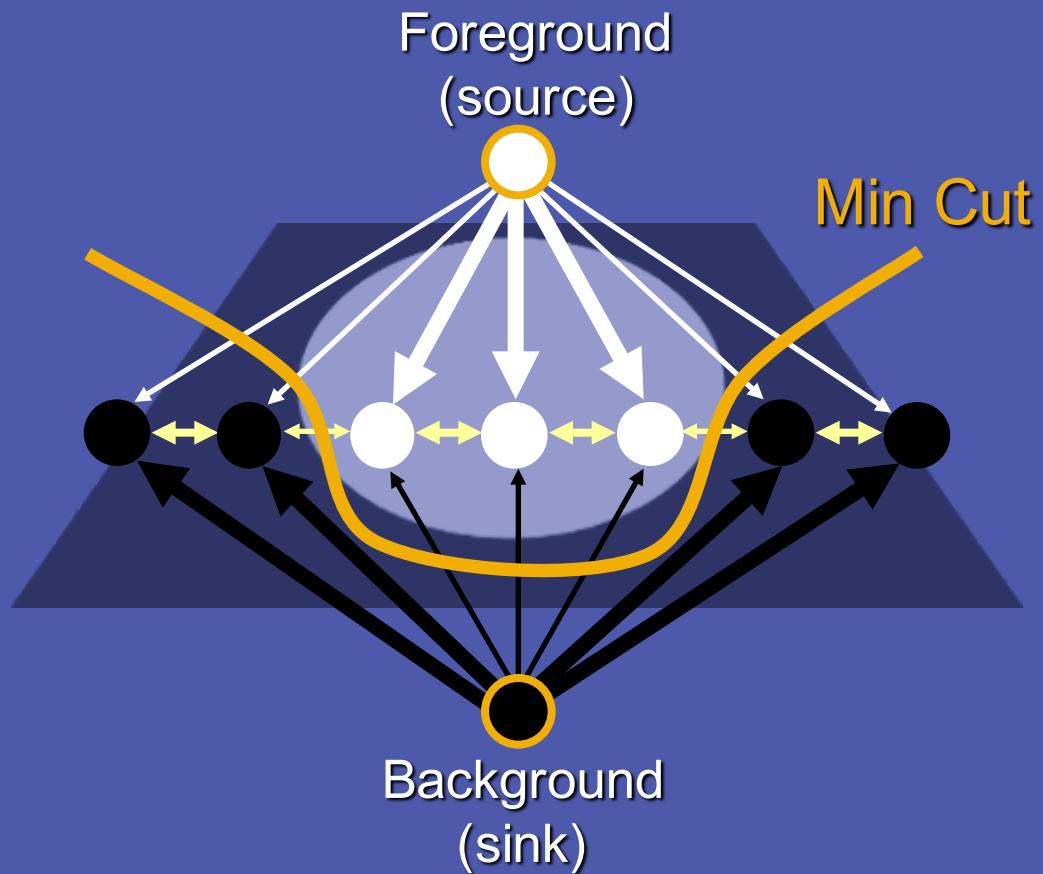
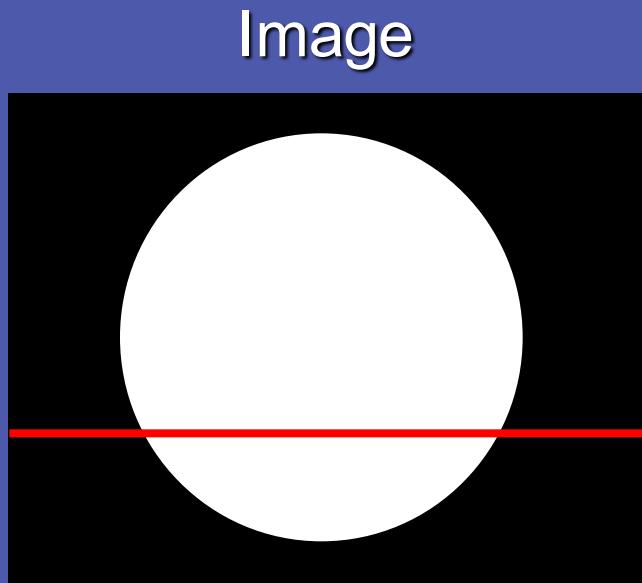
Colour Model



Gaussian Mixture Model (typically 5-8 components)

Graph cuts

Boykov and Jolly (2001)



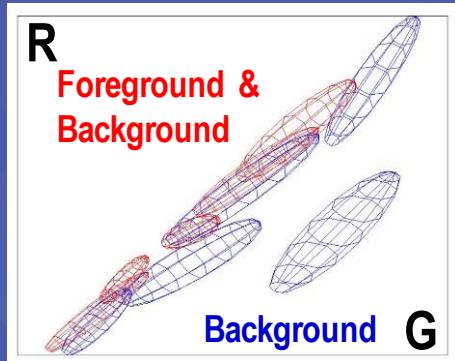
Cut: separating source and sink; Energy: collection of edges

Min Cut: Global minimal energy in polynomial time

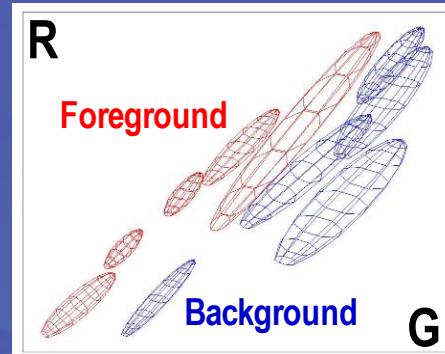
Colour Model



SIGGRAPH2004



Iterated
graph cut



Gaussian Mixture Model (typically 5-8 components)

Source: Rother

GrabCut segmentation

1. Define graph
 - usually 4-connected or 8-connected

2. Define unary potentials
 - Color histogram or mixture of Gaussians for background and foreground

$$\text{unary_potential}(x) = -\log \left(\frac{P(c(x); \theta_{foreground})}{P(c(x); \theta_{background})} \right)$$

3. Define pairwise potentials

$$\text{edge_potential}(x, y) = k_1 + k_2 \exp \left\{ \frac{-\|c(x) - c(y)\|^2}{2\sigma^2} \right\}$$

4. Apply graph cuts

5. Return to 2, using current labels to compute foreground, background models

What is easy or hard about these cases for graphcut-based segmentation?



Easier examples



More difficult Examples

Camouflage &
Low Contrast

Initial
Rectangle



Initial
Result



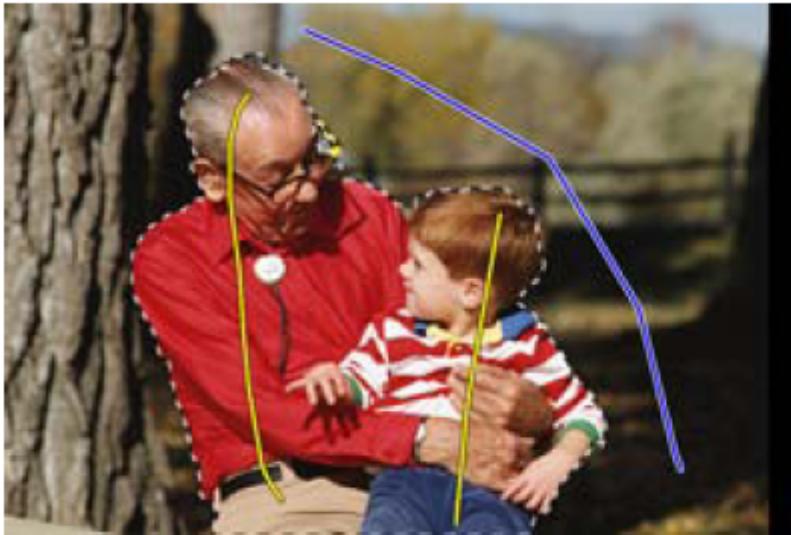
Fine structure



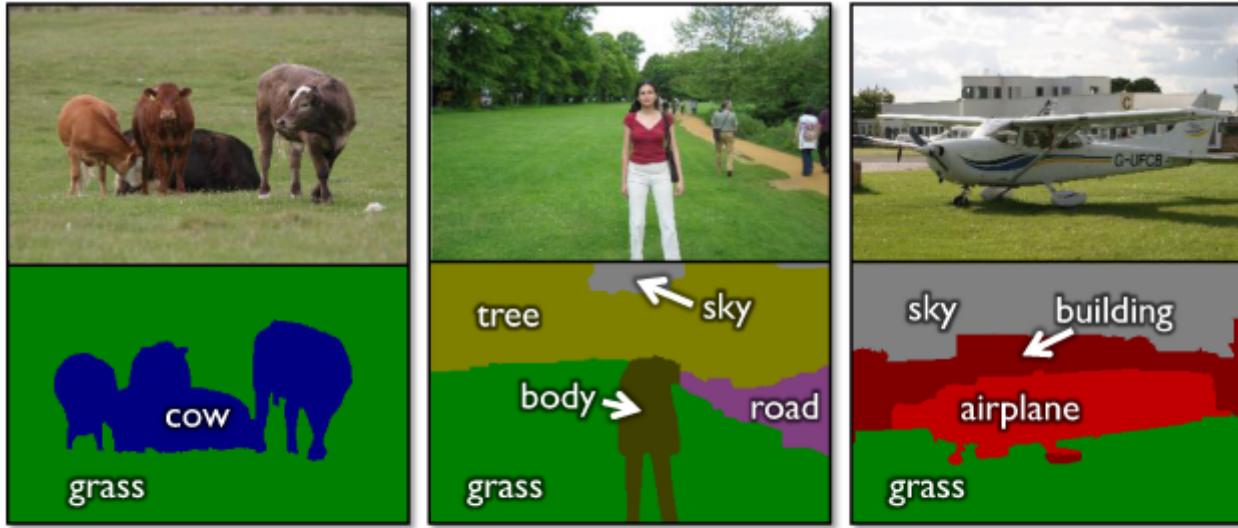
Harder Case



Lazy Snapping (Li et al. SG 2004)

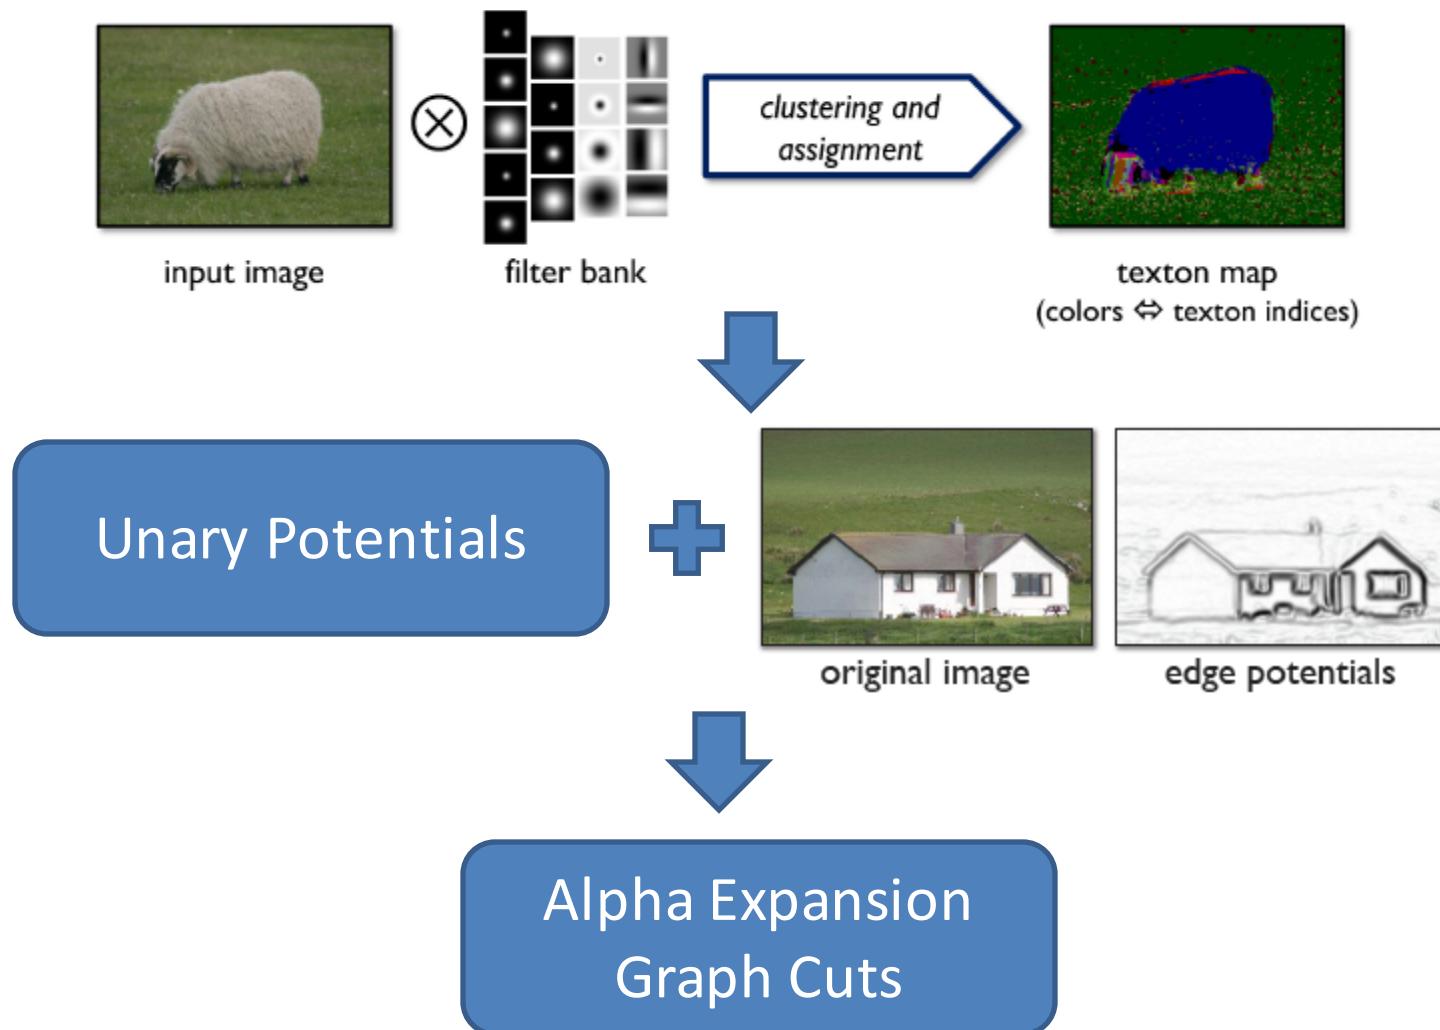


Using graph cuts for recognition



object classes	building	grass	tree	cow	sheep	sky	airplane	water	face	car
bicycle	flower	sign	bird	book	chair	road	cat	dog	body	boat

Using graph cuts for recognition



Limitations of graph cuts

- Associative: edge potentials penalize different labels

Must satisfy

$$E^{i,j}(0,0) + E^{i,j}(1,1) \leq E^{i,j}(0,1) + E^{i,j}(1,0)$$

- If not associative, can sometimes clip potentials
- Approximate for multilabel
 - Alpha-expansion or alpha-beta swaps

Graph cuts: Pros and Cons

- Pros
 - Very fast inference
 - Can incorporate data likelihoods and priors
 - Applies to a wide range of problems (stereo, image labeling, recognition)
- Cons
 - Not always applicable (associative only)
 - Need unary terms (not used for generic segmentation)
- Use whenever applicable

More about MRFs/CRFs

- Other common uses
 - Graph structure on regions
 - Encoding relations between multiple scene elements
- Inference methods
 - Loopy BP or BP-TRW: approximate, slower, but works for more general graphs

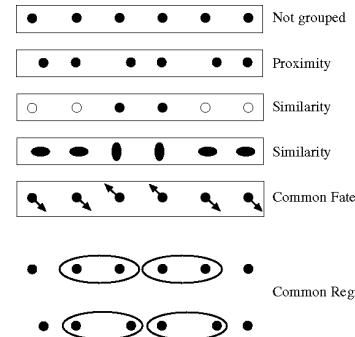
Further reading and resources

- Graph cuts
 - <http://www.cs.cornell.edu/~rdz/graphcuts.html>
 - Classic paper: [What Energy Functions can be Minimized via Graph Cuts?](#) (Kolmogorov and Zabih, ECCV '02/PAMI '04)
- Belief propagation

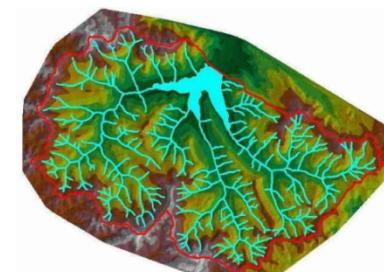
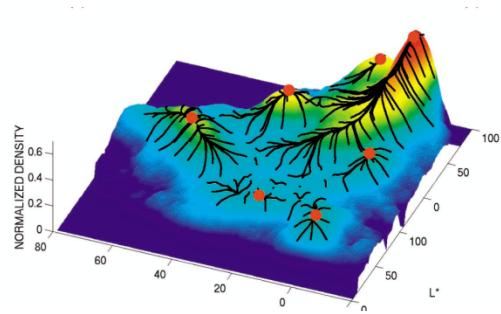
Yedidia, J.S.; Freeman, W.T.; Weiss, Y., "Understanding Belief Propagation and Its Generalizations", Technical Report, 2001:
<http://www.merl.com/publications/TR2001-022/>
- Normalized cuts and image segmentation (Shi and Malik)
<http://www.cs.berkeley.edu/~malik/papers/SM-ncut.pdf>
- N-cut implementation
<http://www.seas.upenn.edu/~timothee/software/ncut/ncut.html>

Next Class

- Gestalt grouping



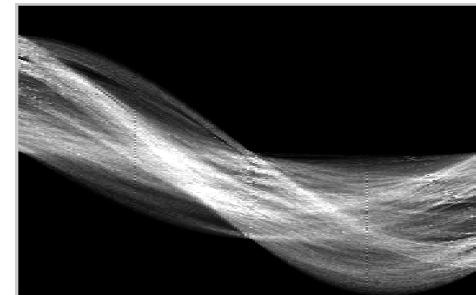
- More segmentation methods



Recap of Grouping and Fitting

Edge and line detection

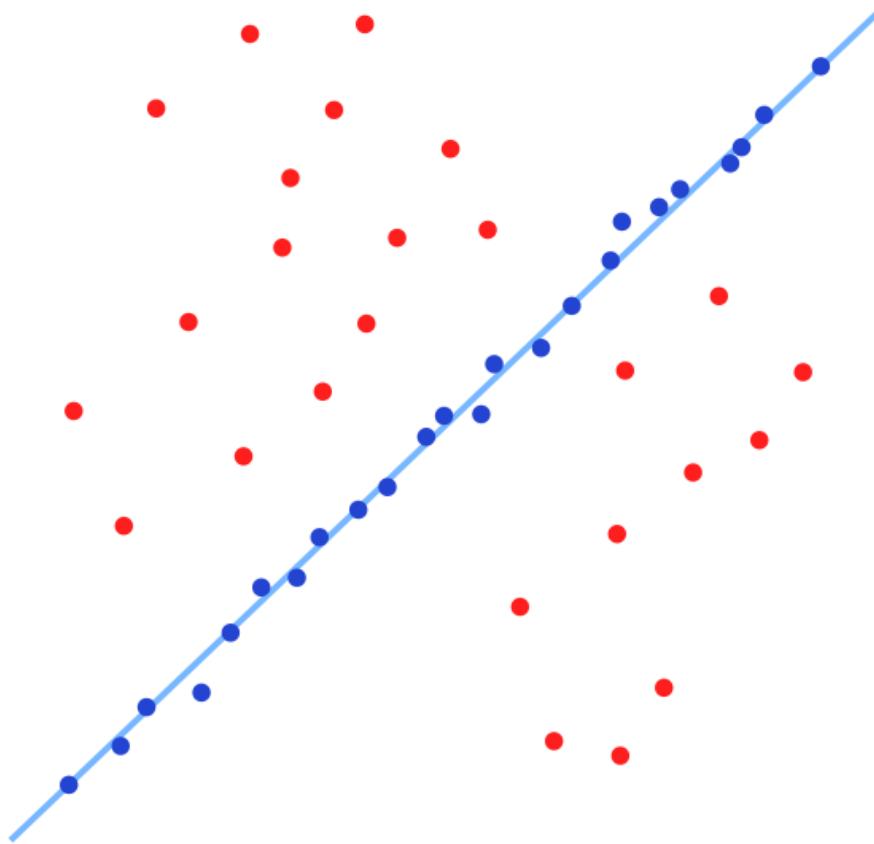
- Canny edge detector = smooth → derivative → thin → threshold → link
- Generalized Hough transform = points vote for shape parameters
- Straight line detector = canny + gradient orientations → orientation binning → linking → check for straightness



Robust fitting and registration

Key algorithms

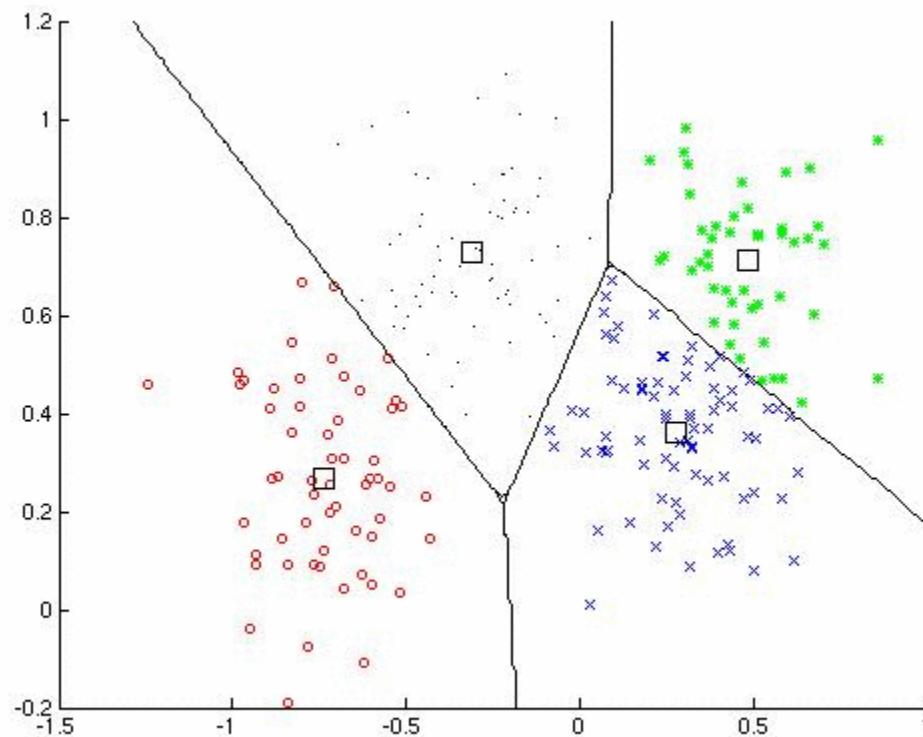
- RANSAC, Hough Transform



Clustering

Key algorithm

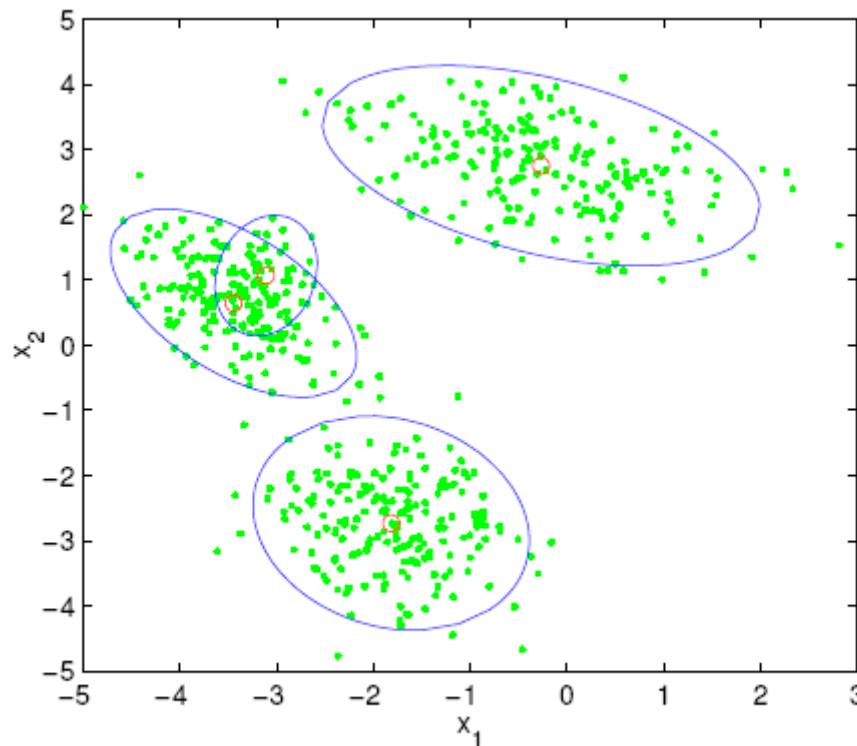
- K-means



EM and Mixture of Gaussians

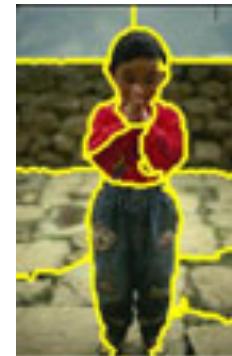
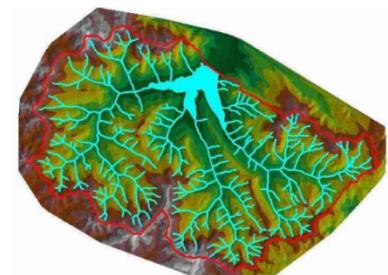
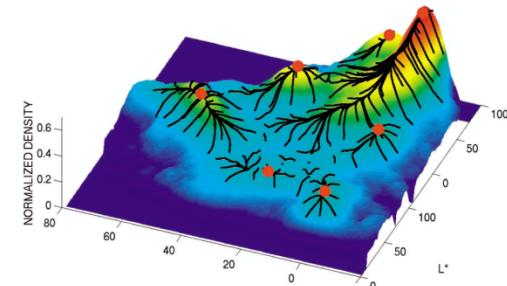
Tutorials:

<http://www.cs.duke.edu/courses/spring04/cps196.1/.../EM/tomasiEM.pdf>
http://www-clmc.usc.edu/~adsouza/notes/mix_gauss.pdf



Segmentation

- Mean-shift segmentation
 - Flexible clustering method, good segmentation
- Watershed segmentation
 - Hierarchical segmentation from soft boundaries
- Normalized cuts
 - Produces regular regions
 - Slow but good for oversegmentation
- MRFs with Graph Cut
 - Incorporates foreground/background/object model and prefers to cut at image boundaries
 - Good for interactive segmentation or recognition



Next section: Recognition

- How to recognize
 - Specific object instances
 - Faces
 - Scenes
 - Object categories